

Module-4

MEMORY

Memory Organization.			
4.0	4.1	Characteristics of memory system and hierarchy, Main memory, Cache memory principles , Elements of Cache Design	
	4.2	ROM, Types of ROM, RAM, SRAM, DRAM, Flash memory, High speed memories	11 CO3
	4.3	Cache Memory Organization: Address mapping, Replacement Algorithms, Cache Coherence, MESI protocol, Interleaved and associative memories, Virtual memory, Main memory allocation, Segmentation ,Paging, Secondary storage, RAID levels	

Characteristics of Memory

- 1.** Location
- 2.** Capacity
- 3.** Unit of transfer
- 4.** Access method
- 5.** Performance(SRAM,DRAM)
- 6.** Physical type
- 7.** Physical characteristics
- 8.** Organisation
 - 1.** Direct Mapping,Associative Mapping

1. Location

- CPU
- Internal
- External



Memory Card Reader



USB Flash Memory



Media Devices



External Optical Drives



ZIP Drive

2. Capacity

- Word size
 - The natural unit of organisation
- Number of words
 - or Bytes

Types of various Units of Memory-

Byte

Kilo Byte

Mega Byte

Giga Byte

Tera Byte

Peta Byte

Exa Byte

Zetta Byte

Yotta Byte

NAME	EQUAL TO	SIZE(IN BYTES)
Bit	1 bit	1/8
Nibble	4 bits	1/2 (rare)
Byte	8 bits	1
Kilobyte	1024 bytes	1024
Megabyte	1, 024kilobytes	1, 048, 576
Gigabyte	1, 024 megabytes	1, 073, 741, 824
Terrabyte	1, 024 gigabytes	1, 099, 511, 627, 776
Petabyte	1, 024 terrabytes	1, 125, 899, 906, 842, 624
Exabyte	1, 024 petabytes	1, 152, 921, 504, 606, 846, 976
Zettabyte	1, 024 exabytes	1, 180, 591, 620, 717, 411, 303, 424
Yottabyte	1. 024 zettabytes	1. 208. 925. 819. 614. 629. 174. 706. 176

3. Unit of Transfer

- **INTERNAL**
 - Usually governed by **data bus width**
- **EXTERNAL**
 - Usually a **block** which is much larger than a word
- **Addressable unit**
 - Smallest location which can be uniquely addressed

4. Access Methods (1)

- **Sequential**

- Start at the beginning and read through in order
- Access time depends on location of data and previous location
- e.g. tape



- **Direct**

- Individual blocks have unique address
- Access is by jumping to vicinity plus sequential search
- Access time depends on location and previous location
- e.g. disk



Access Methods (2)

- **Random**

- Individual addresses identify locations exactly
- Access time is independent of location or previous access
- e.g. RAM

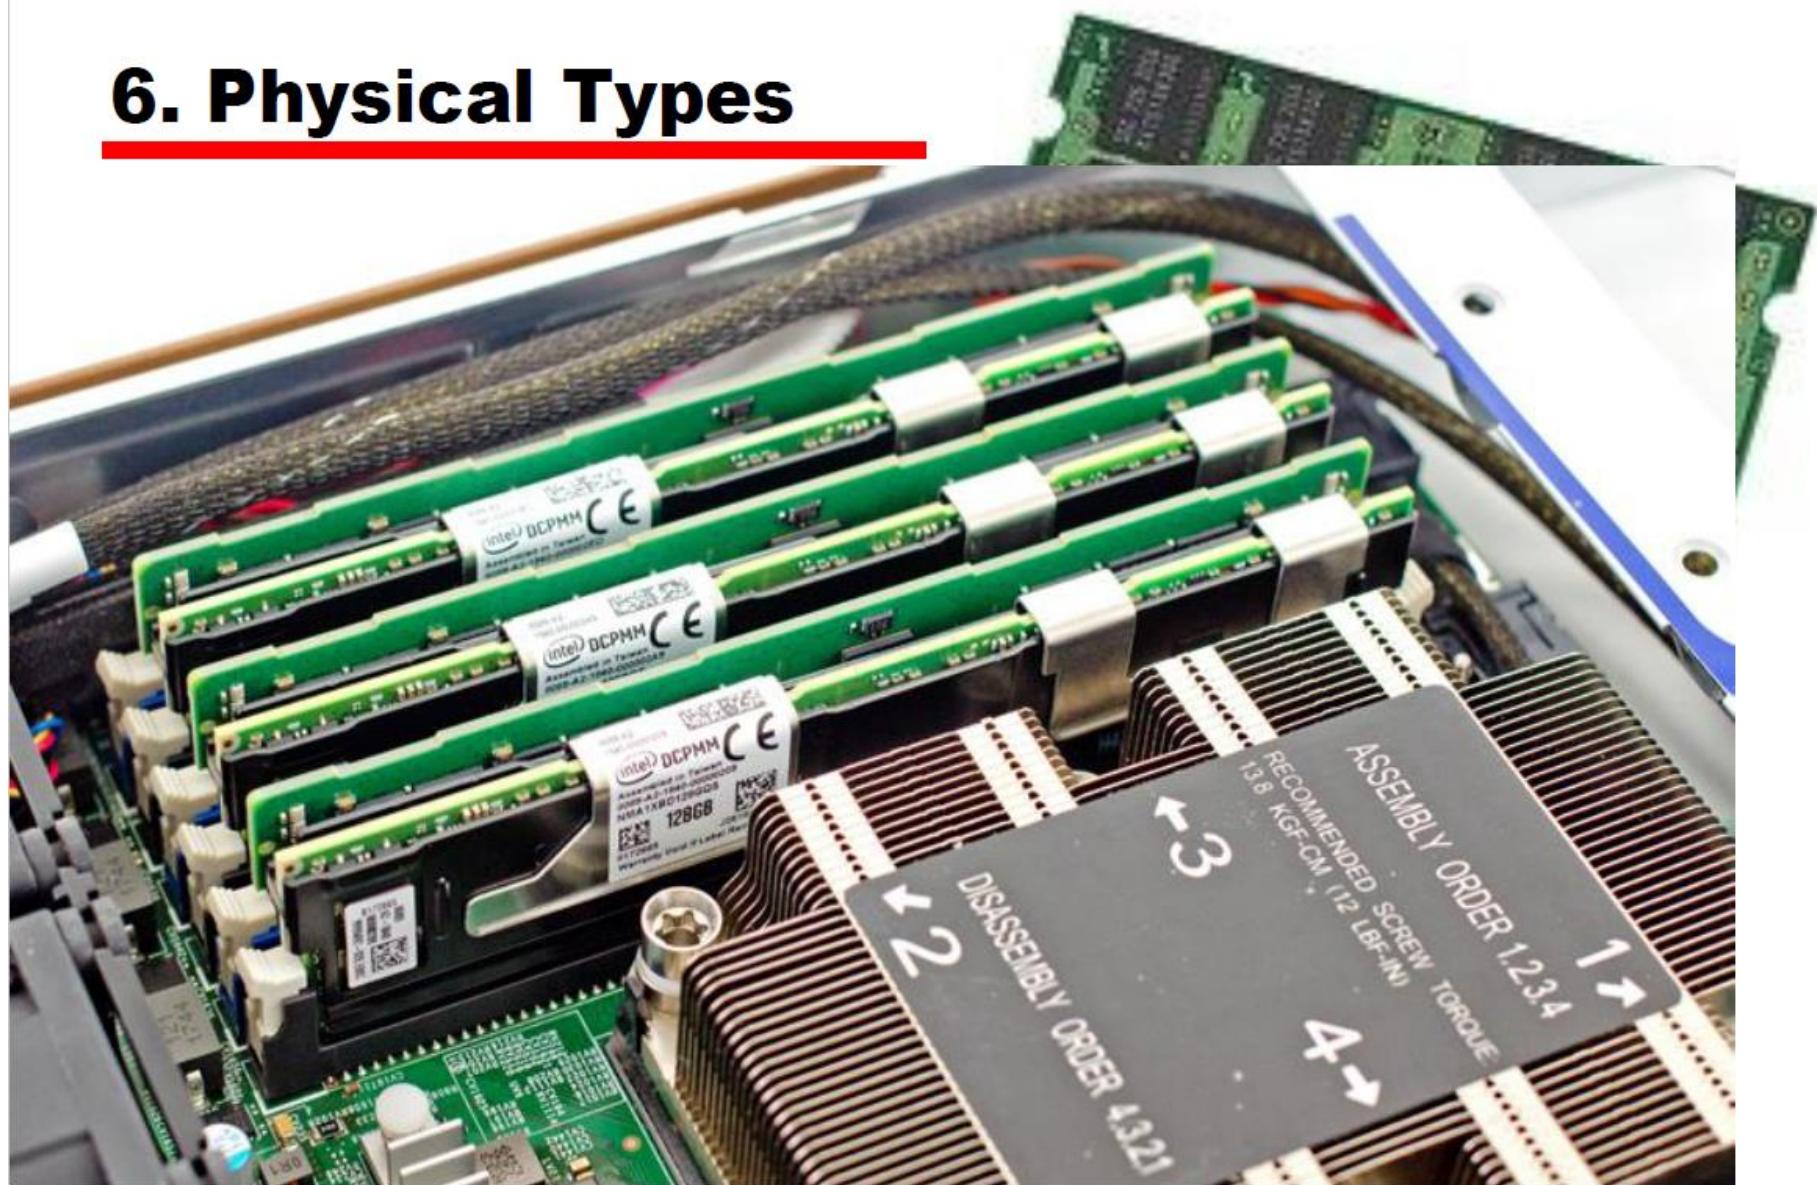
- **Associative**

- Data is located by a comparison with contents of a portion of the store
- Access time is independent of location or previous access
- e.g. cache

5. Performance

- **Access time**
 - Time between requesting for operation and the time it is made available at the required location
- **Memory Cycle time**
 - Minimum time elapsed between two **consecutive read requests**
- **Transfer Rate**
 - Rate at which data can be moved

6. Physical Types



PHYSICAL TYPES

- **physical types-**
- semiconductor memory,
- magnetic surface memory,-used for disk and tape,
- optical and magneto-optical

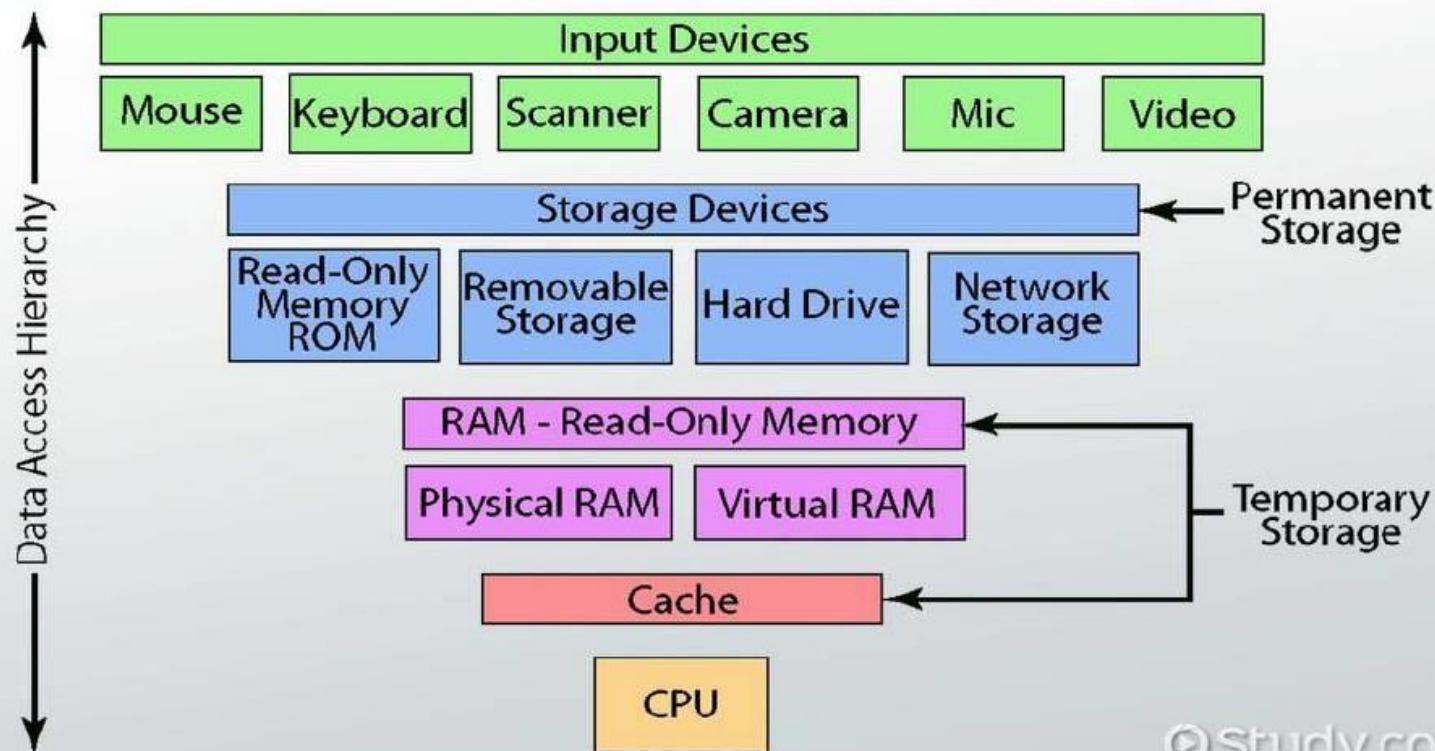
7. Physical Characteristics

- Decay
- Volatility
- Erasable
- Power consumption

8. Organisation

- Physical arrangement of bits into words
- Not always obvious
- e.g. interleaved

HOW A COMPUTER USES MEMORY

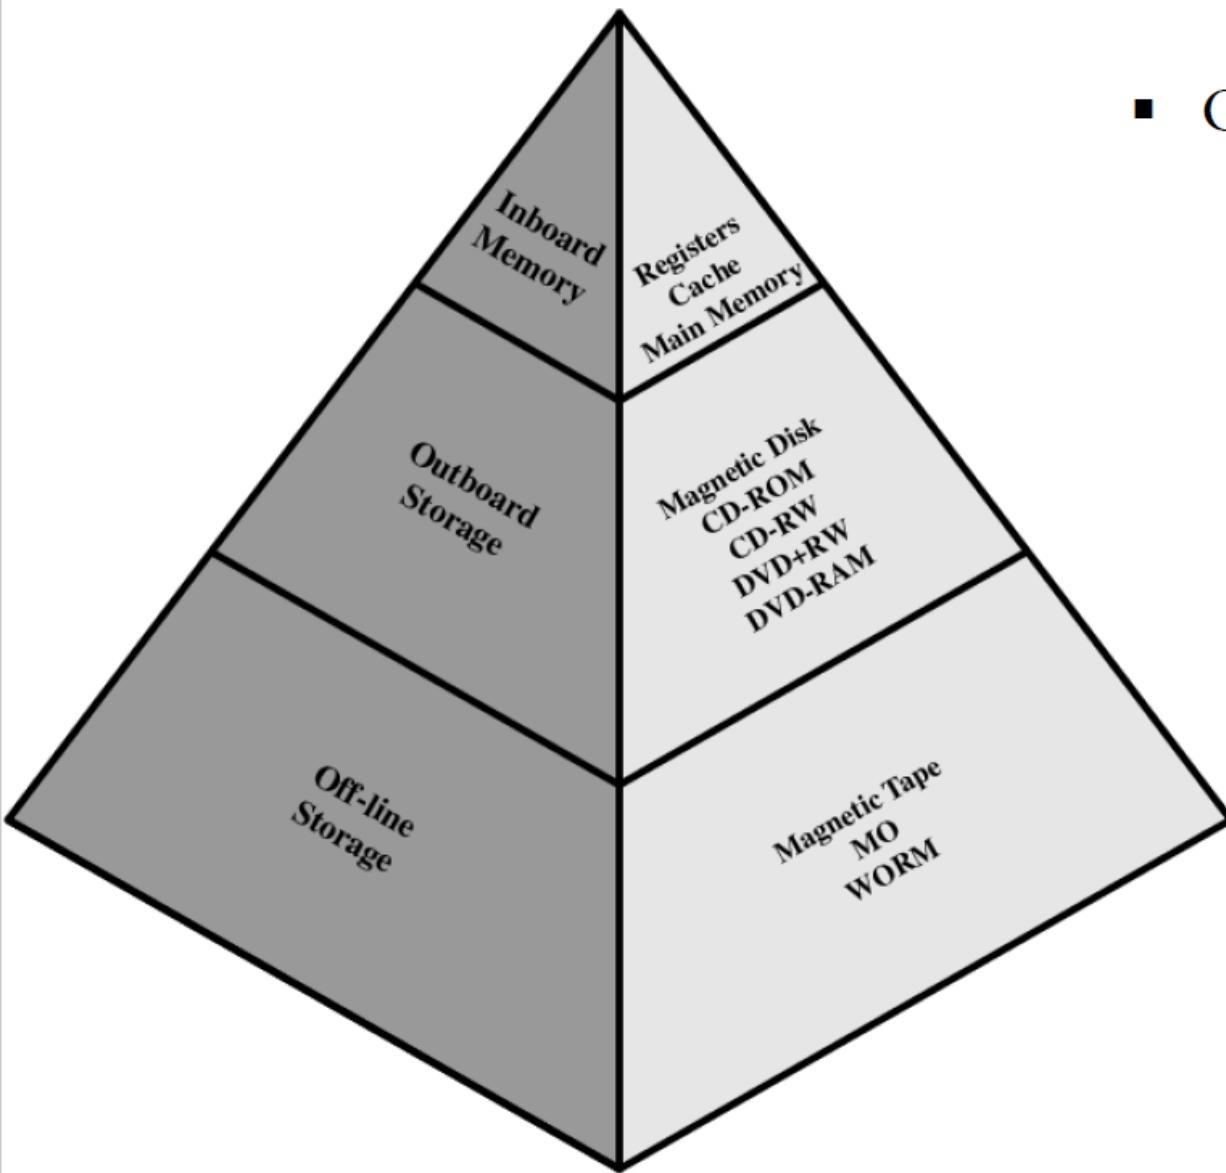


Location	Internal (e.g. processor registers, main memory, cache)	Performance	Access time
	External (e.g. optical disks, magnetic disks, tapes)	Cycle time	Transfer rate
Capacity		Physical Type	
Number of words		Semiconductor	
Number of bytes		Magnetic	
Unit of Transfer		Optical	
Word		Magneto-optical	
Block		Physical Characteristics	
Access Method		Volatile/nonvolatile	
Sequential		Erasable/nonerasable	
Direct		Organization	
Random		Memory modules	
Associative			

Memory Hierarchy

- Registers
 - In CPU
- Internal or Main memory
 - May include one or more levels of cache
 - “RAM”
- External memory
 - Backing store

Memory Hierarchy - Diagram



- Going down the hierarchy
 - Decreasing **Cost**
 - Increase **Capacity**
 - Increase **Access Time**

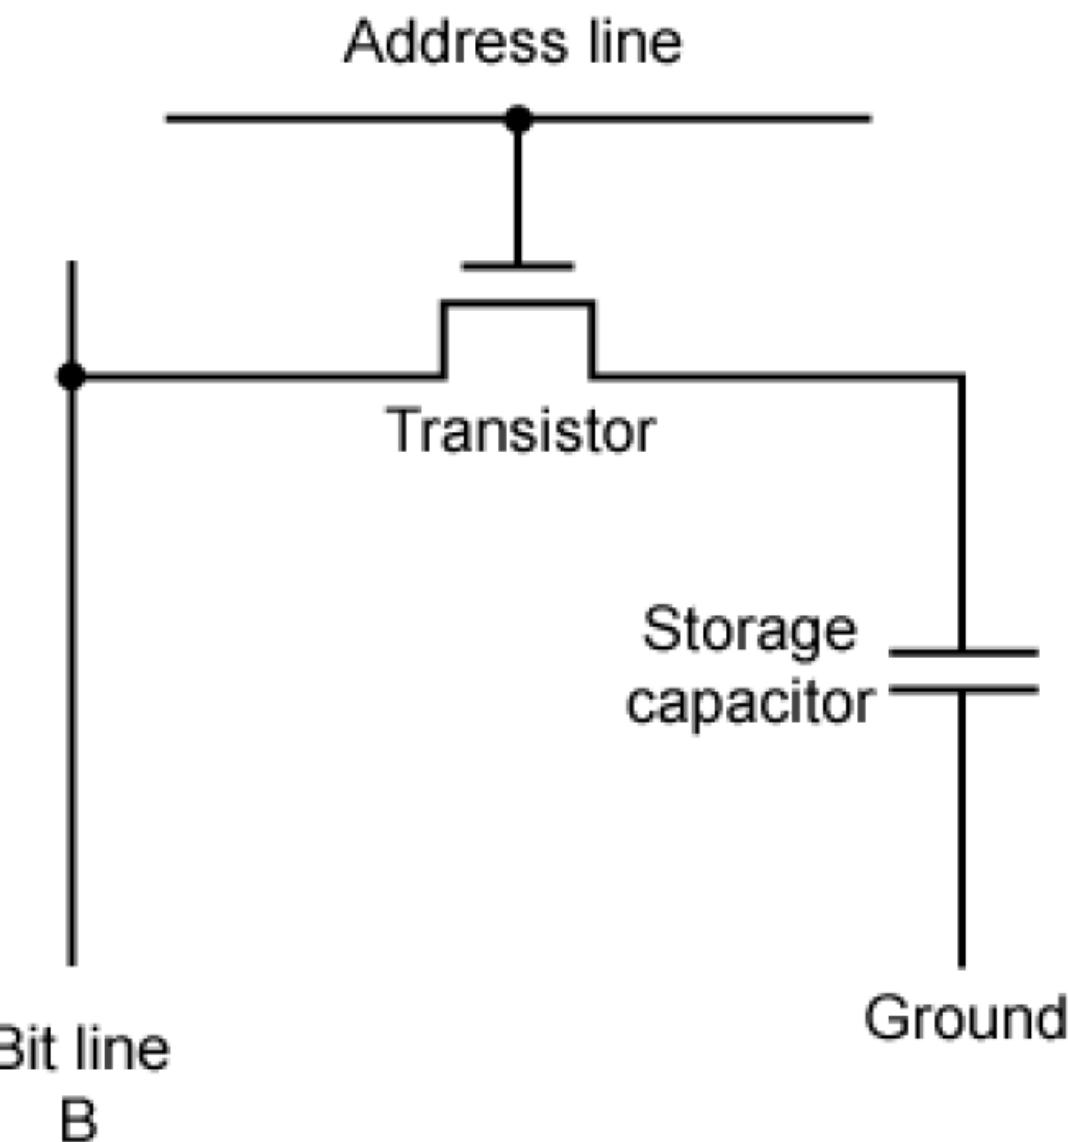
RAM- Random Access Memory

- RAM
 - random access
 - Read/Write
 - Volatile
 - Temporary storage
 - Static or dynamic

Dynamic RAM

- Bits stored as charge in capacitors
- Charges leak
- Need **refreshing** even when powered
- Simpler construction
- Less expensive
- Need refresh circuits
- Slower
- Main memory
- Essentially analog
 - Level of charge determines value

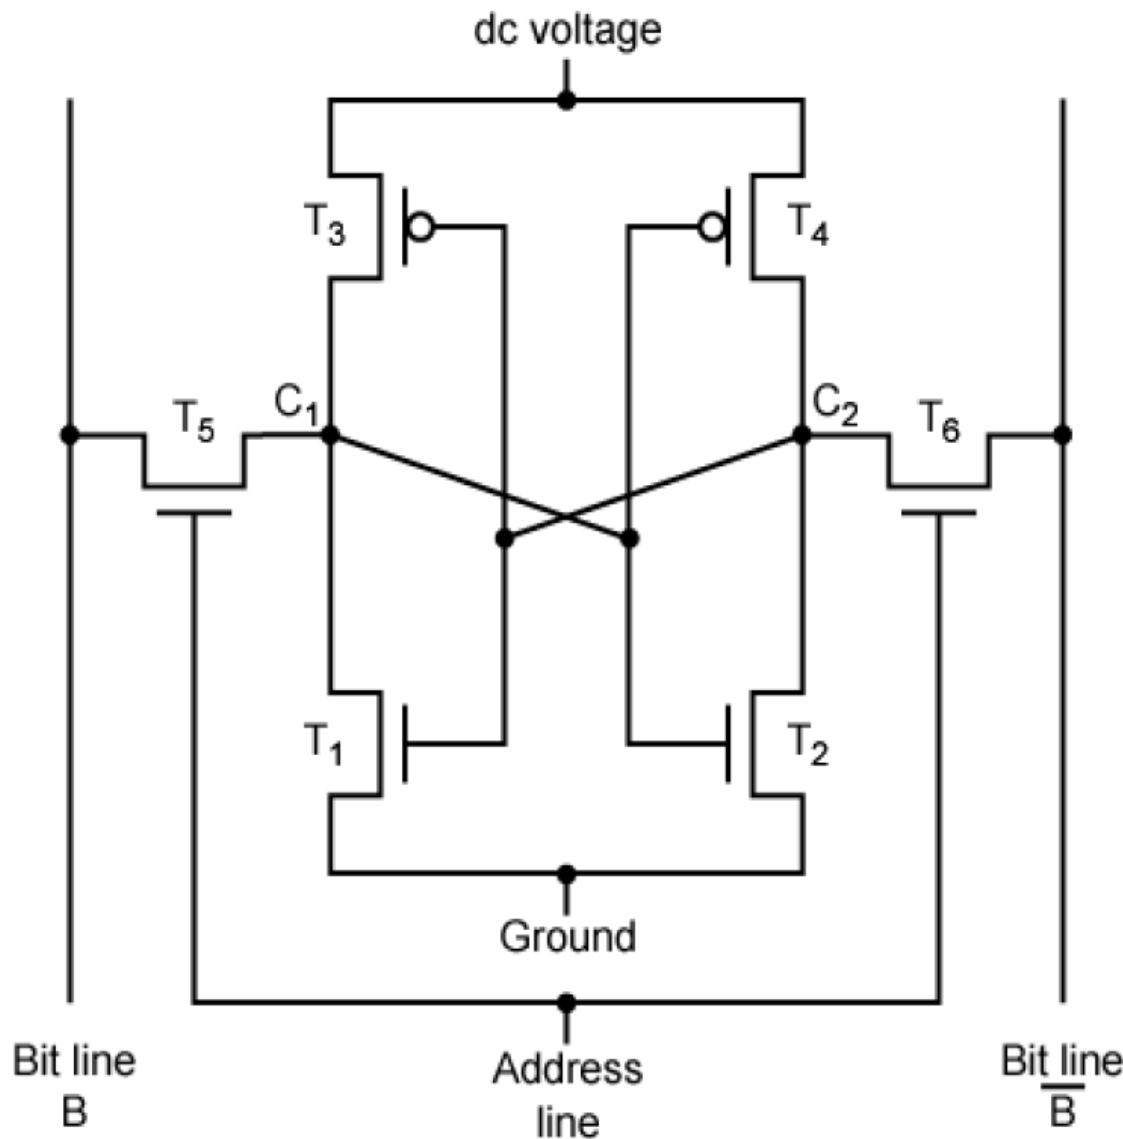
Dynamic RAM Structure



Static RAM

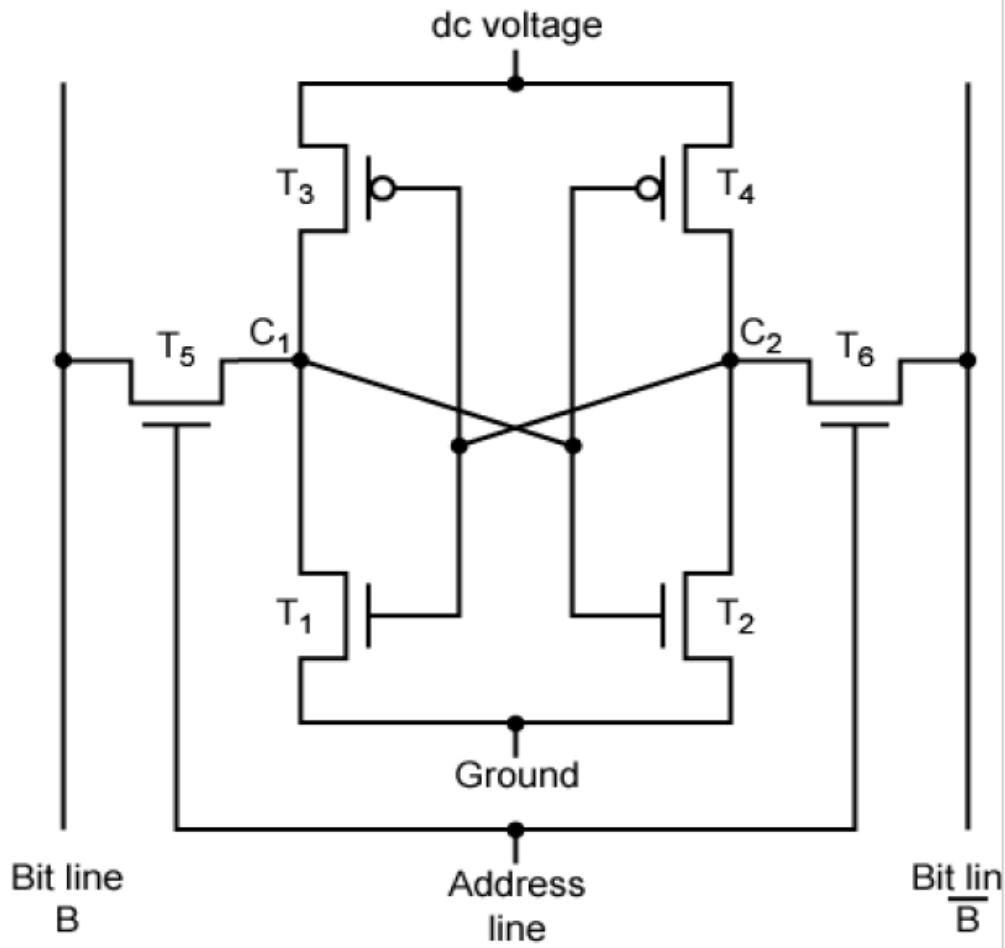
- Bits stored as on/off switches
- No charges to leak
- No refreshing needed when powered
- More complex construction
- Larger information per bit
- More expensive
- Does not need refresh circuits
- Faster
- Cache
- Digital
 - Uses flip-flops

Static RAM Structure



Static RAM Operation

- Transistor arrangement gives stable logic state
- State 1
 - C_1 high, C_2 low
 - $T_1 T_4$ off, $T_2 T_3$ on
- State 0
 - C_2 high, C_1 low
 - $T_2 T_3$ off, $T_1 T_4$ on



SRAM v DRAM

- Both volatile
 - Power needed to preserve data
- Dynamic cell
 - Simpler to build, smaller
 - More dense
 - Less expensive
 - Needs refresh
 - Larger memory units
- Static
 - Faster
 - Cache

Read Only Memory (ROM)

- Permanent storage
 - Non volatile
 - Can read a ROM but can't write new data into it
- TYPES OF ROM
 - PROM
 - EPROM
 - EEPROM

PROM-Programmable Read Only Memory

- **PROM-Programmable ROM**
- •Written by user.
- •Programmable (“once”)
- —**PROM**
- —Small amount of data to be written
- —Less expensive
- —Non volatile, written only once
- —Writing performed electrically by the user(fabricated without writing)

Read “mostly”

- Erasable Programmable (**EPROM**)
 - Erased by UV
- Electrically Erasable (**EEPROM**)
 - Takes much longer to write than read
- Flash memory
 - Erase whole memory electrically

EPROM

- Read and write-
- —All storage cells should be erased to **initial state** by exposure to UV radiation .
- —Can be altered multiple times and holds data virtually indefinitely
- —More expensive than PROM

EEPROM

- Can be written anytime without erasing prior contents
- Write operation takes longer than read
- More expensive than EPROM, less dense

Types of ROM

1. Programmable Read Only Memory (PROM)

- Empty of data when manufactured
- May be permanently programmed by the user

2. Erasable Programmable Read Only Memory (EPROM)

- Can be programmed, erased and reprogrammed
- The EPROM chip has a small window on top allowing it to be erased by shining ultra-violet light on it
- After reprogramming the window is covered to prevent new contents being erased
- Access time is around 45 - 90 nanoseconds

Types of ROM

3. Electrically Erasable Programmable Read Only Memory (EEPROM)

- Reprogrammed electrically without using ultraviolet light
- Must be removed from the computer and placed in a special machine to do this
- Access times between 45 and 200 nanoseconds

4. Flash ROM

- Similar to EEPROM
- However, can be reprogrammed while still in the computer
- Easier to upgrade programs stored in Flash ROM
- Used to store programs in devices e.g. modems
- Access time is around 45 - 90 nanoseconds

5. ROM cartridges

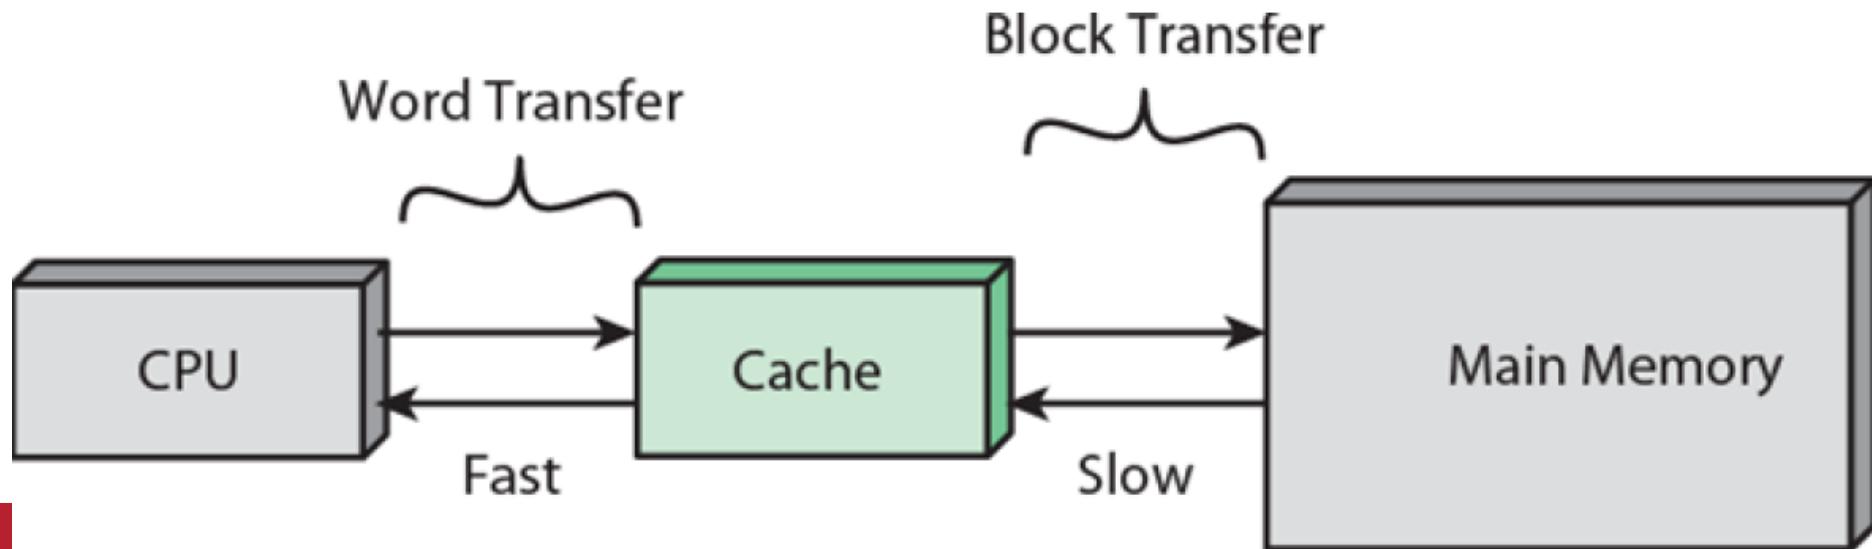
- Commonly used in games machines
- Prevents software from being easily copied

CACHE MEMORY

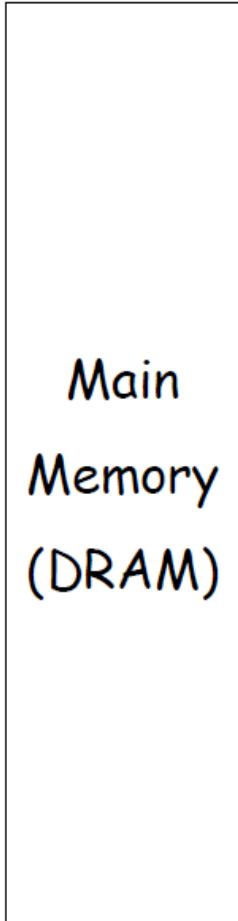
- Cache memory is intended to give –
- memory speed approaching that of the fastest memories available.
- Provide a large memory size at the price of less expensive types of semiconductor memories.
- A relatively large and slow main memory together with a
- smaller, faster cache memory.
- The cache contains a copy of portions of main memory.

Cache

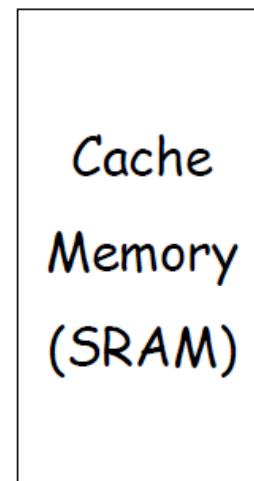
- **Small** amount of fast memory
- Sits between main memory and CPU
- May be located on CPU chip or module



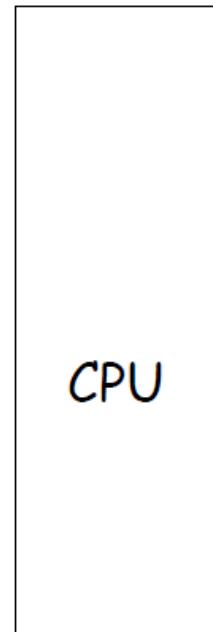
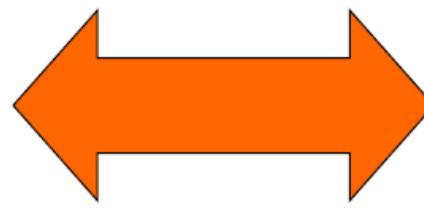
The operation of cache memory



1 Cache fetches data from addresses in main memory



2. CPU checks to see whether the next instruction it requires is in cache



4. If not, the CPU has to fetch next instruction from main memory - a much slower process

3. If present then the instruction is fetched from the cache - a very fast operation

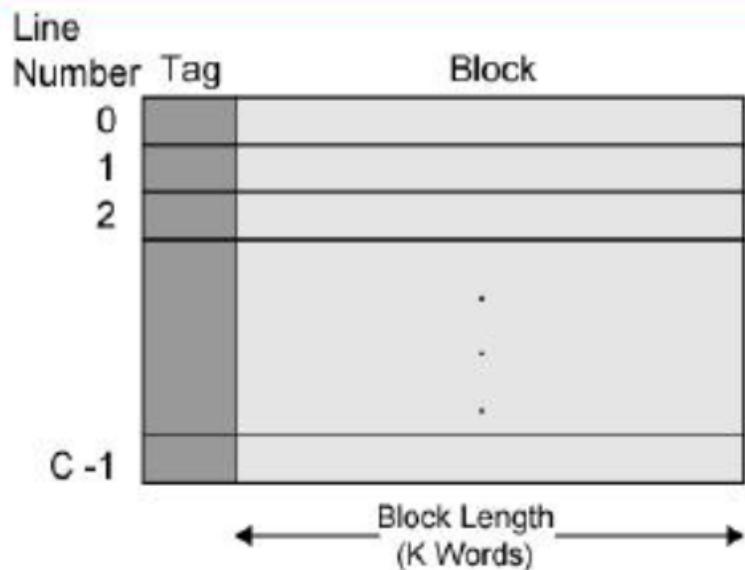


= Bus connections

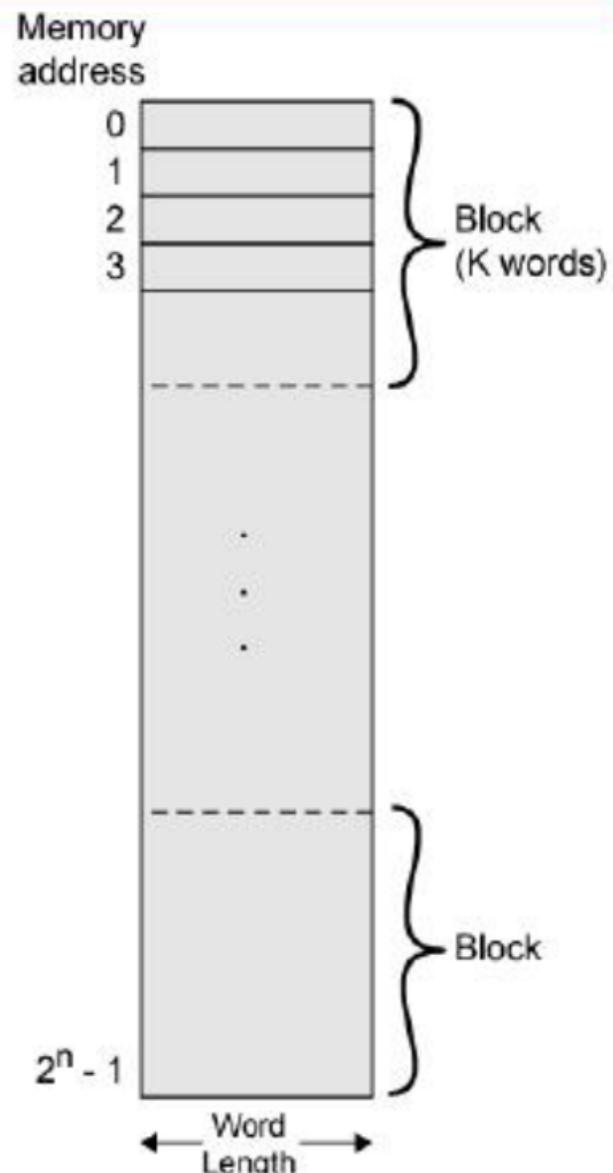
Cache operation – overview

- CPU requests contents of memory location
- Check cache for this data
- If present, get from cache (fast)
- If not present, read required block from main memory to cache
- Then deliver from cache to CPU
- Cache includes **tags** to identify which block of main memory is in each cache slot

Cache/Main Memory Structure



(a) Cache



TAG - A unique identifier for a group of data.

Since different regions of memory may be mapped into a block, the tag is used to differentiate between them.

- CPU cache is divided into three main 'Levels', L1, L2, and L3.
 - The hierarchy is according to the speed, and thus, the size of the cache.
- **Level 1 (L1) cache** -fast small, **embedded** in the processor chip (CPU).
- **Level 2 (L2) cache** more capacity than L1; **located** on the CPU or on a separate chip or coprocessor.
- **Level 3 (L3) cache** specialized memory to improve the performance of L1 and L2.
 - Slower than L1 or L2, double the speed of RAM

Size does matter

- Cost
 - More cache is expensive
- Speed
 - More cache is faster (up to a point)
 - Checking cache for data takes time

Cache Design

- Size
- Mapping Function
- Replacement Algorithm
- Write Policy
- Block Size
- Number of Caches

- Because there are fewer cache lines than main memory blocks, an algorithm is needed for mapping main memory blocks into cache lines.
- Also, a means is needed for determining which main memory block currently occupies a cache line.
- The choice of the mapping function decides how the cache is organized.
- Three techniques can be used: direct, associative, and set associative.

MAPPING TECHNIQUES

- **DIRECT MAPPING**
- **ASSOCIATIVE MAPPING**
 - FULLY ASSOCIATIVE MAPPING**
 - SET ASSOCIATIVE MAPPING**
 - **2-WAY SET ASSOCIATIVE MAPPING**

DIRECT MAPPING CONCEPT

- **Direct Mapping:** Each block from main memory has only one possible place in the cache organization in this technique.
- As in the diagram shown, one block of Main memory to one particular block(line) of cache.

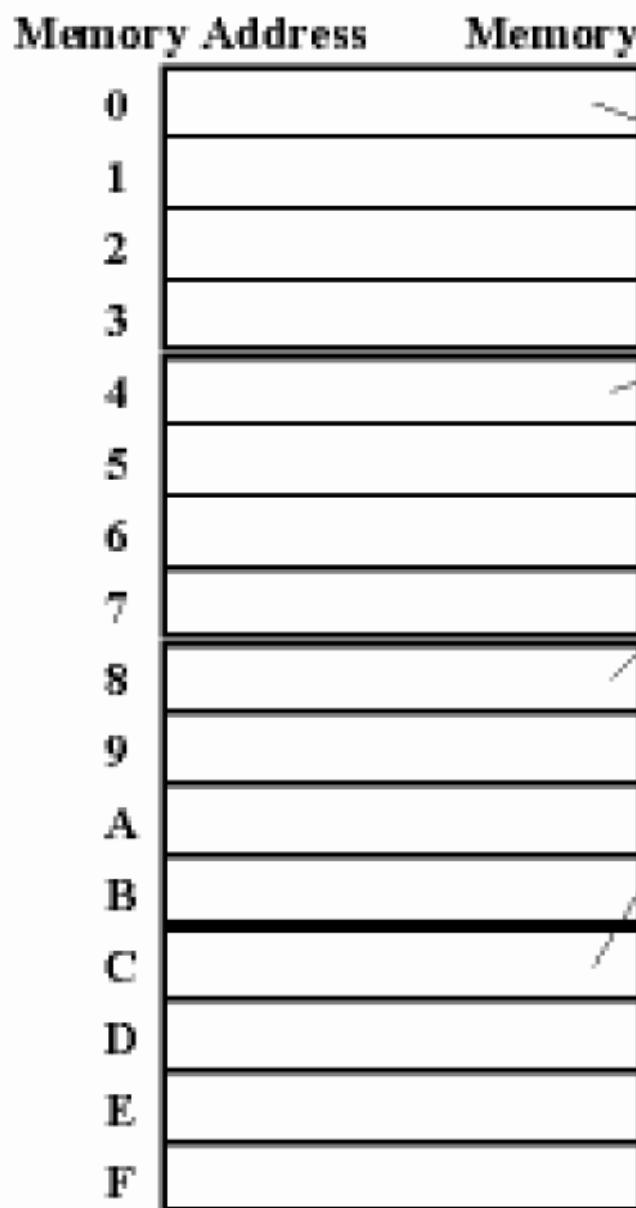
Main Memory

0ABCCE
1
2
3
4FFFFE
5
6
7
8
9
10
11

Cache

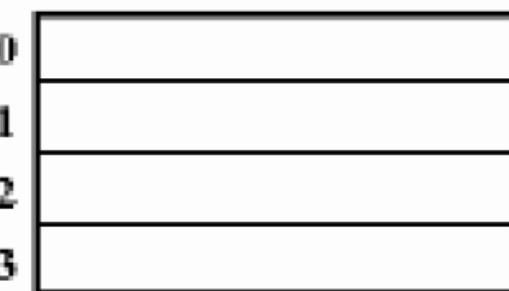
0 / 4 / 8 [4FFFFEE]
1 / 5 / 9 [9]
2 / 6 / 10 [2]
3 / 7 / 11 [7]

DIRECT MAPPING CONCEPT



4 Byte Direct Mapped Cache

Cache Index



K J Somaiya College of Engineering

Explaining with Example

- Consider an MM of 64 bits , block size=4 words
- No : of blocks in MM= $64/4=16$

- Cache =16 words
- Block size=4 words.

O

0

1

2

3

1

4

5

6

7

2

8

9

10

11

3

12

13

14

15

4

16

17

18

19

5

20

21

22

23

6

24

25

26

27

7

28

29

30

31

8

32

33

34

35

9

36

37

38

39

10

40

41

42

43

11

44

45

46

47

12

48

49

50

51

13

52

53

54

55

14

56

57

58

59

15

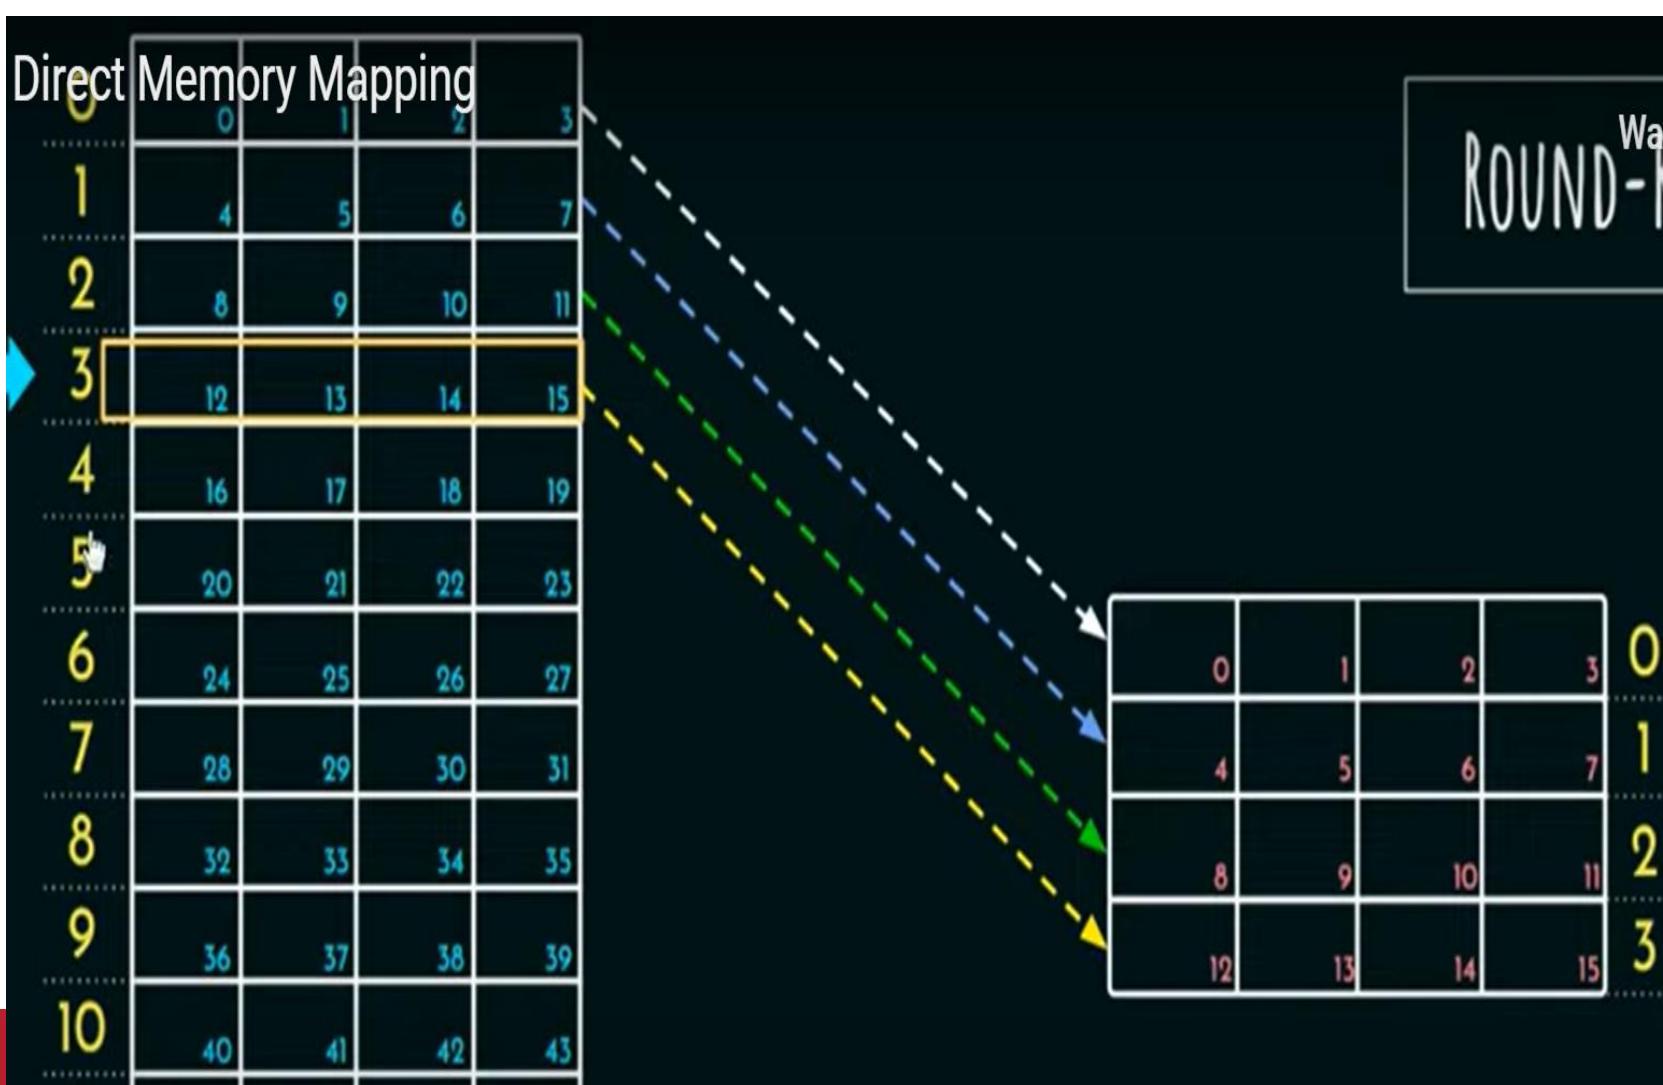
60

61

62

63

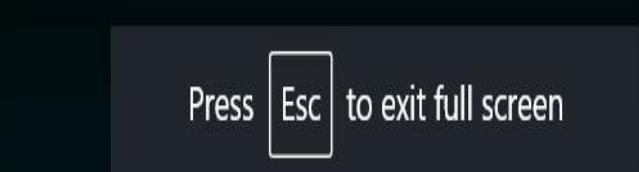
Mapping



Correct Memory Mapping

Press Esc to exit full screen

0	1	2	3
4	5	6	7
8	9	10	11
12	13	14	15
16	17	18	19
20	21	22	23
24	25	26	27
28	29	30	31
32	33	34	35
36	37	38	39
40	41	42	43
44	45	46	47
48	49	50	51
52	53	54	55
56	57	58	59
60	61	62	63
64	65	66	67
68	69	70	71
72	73	74	75
76	77	78	79
80	81	82	83
84	85	86	87
88	89	90	91
92	93	94	95
96	97	98	99



16	17	18	19
20	21	22	23
8	9	10	11
12	13	14	15

0
1
2
3

DIRECT MAPPING

- Cache- 128 blocks of 16 words each
 - $128 * 16 = 2048$ **approx 2 KB**

0	16 words
1	16 words
...	16 words
127	16 words

- Main Memory – 4K blocks of 16 words each
 - $4K * 16 = 64000$ **approx 64 KB**
- **TOTAL ADDRESS SIZE 16 bit**

Tag	Block/Line(128)	Word(16)
-----	-----------------	----------

5($16-(7+4)$)

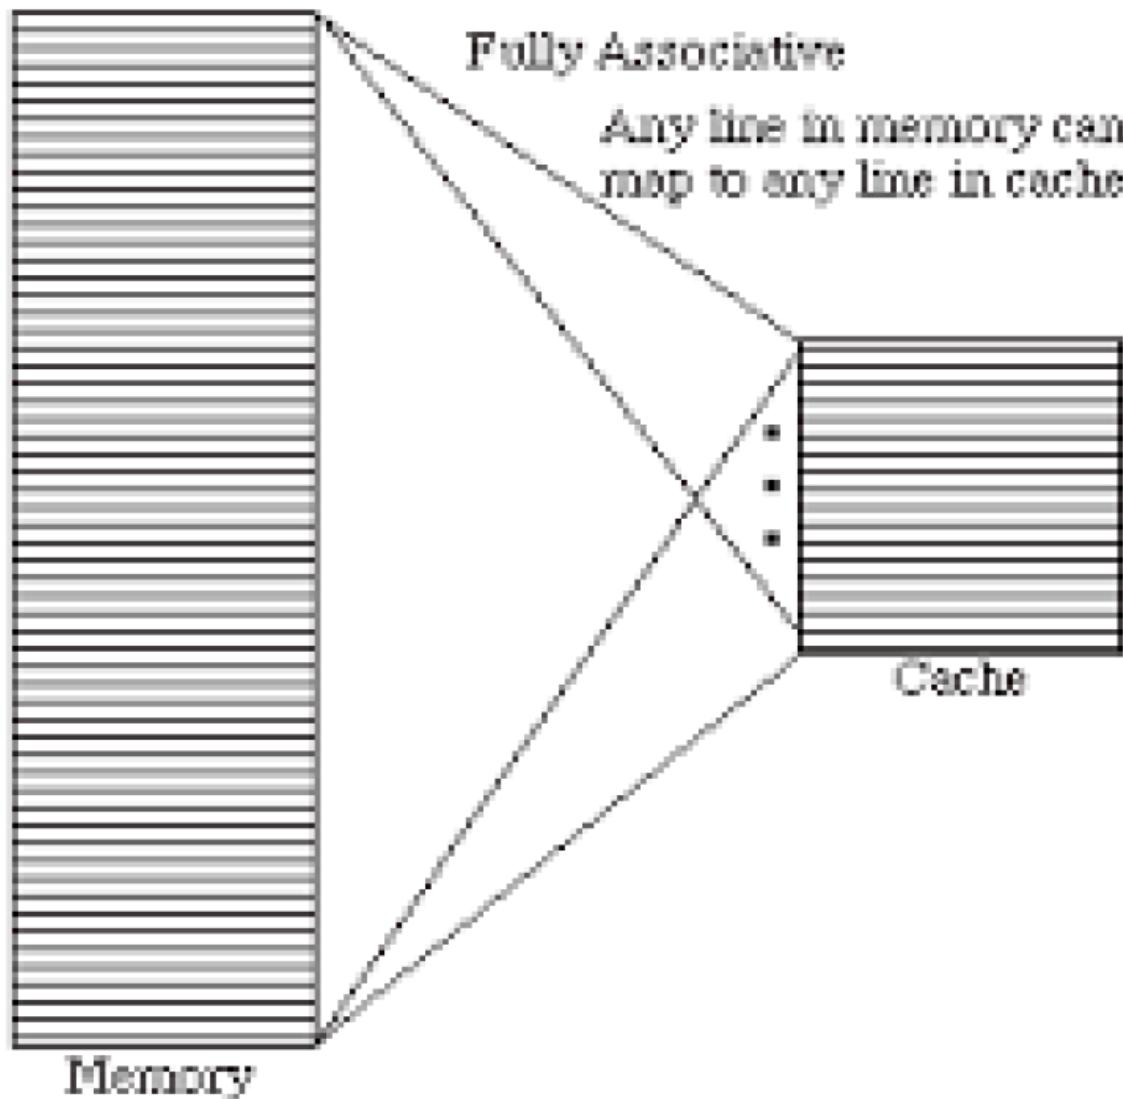
7(2^7)

4 (2^4)

Working of Direct Mapping

- "Word" field selects one from among the 16 addressable words in a line.
- The "Line" field defines the cache line where this memory line should reside.
- The "Tag" field of the address is then compared with that cache line's 5-bit tag to determine whether there is a hit or a miss.
 - If there's a miss, we need to swap out the memory line that occupies that position in the cache and replace it with the desired memory line.

FULLY ASSOCIATIVE MAPPING



ASSOCIATIVE MAPPING

- Cache- 128 blocks of 16 words each
- $128 * 16 = 2048$ **approx 2 KB**
- Main Memory – 4K blocks of 16 words each
- $4K * 16 = 64000$ **approx 64 KB**
- **TOTAL ADDRESS SIZE 16 bit**

Tag	Word
-----	------

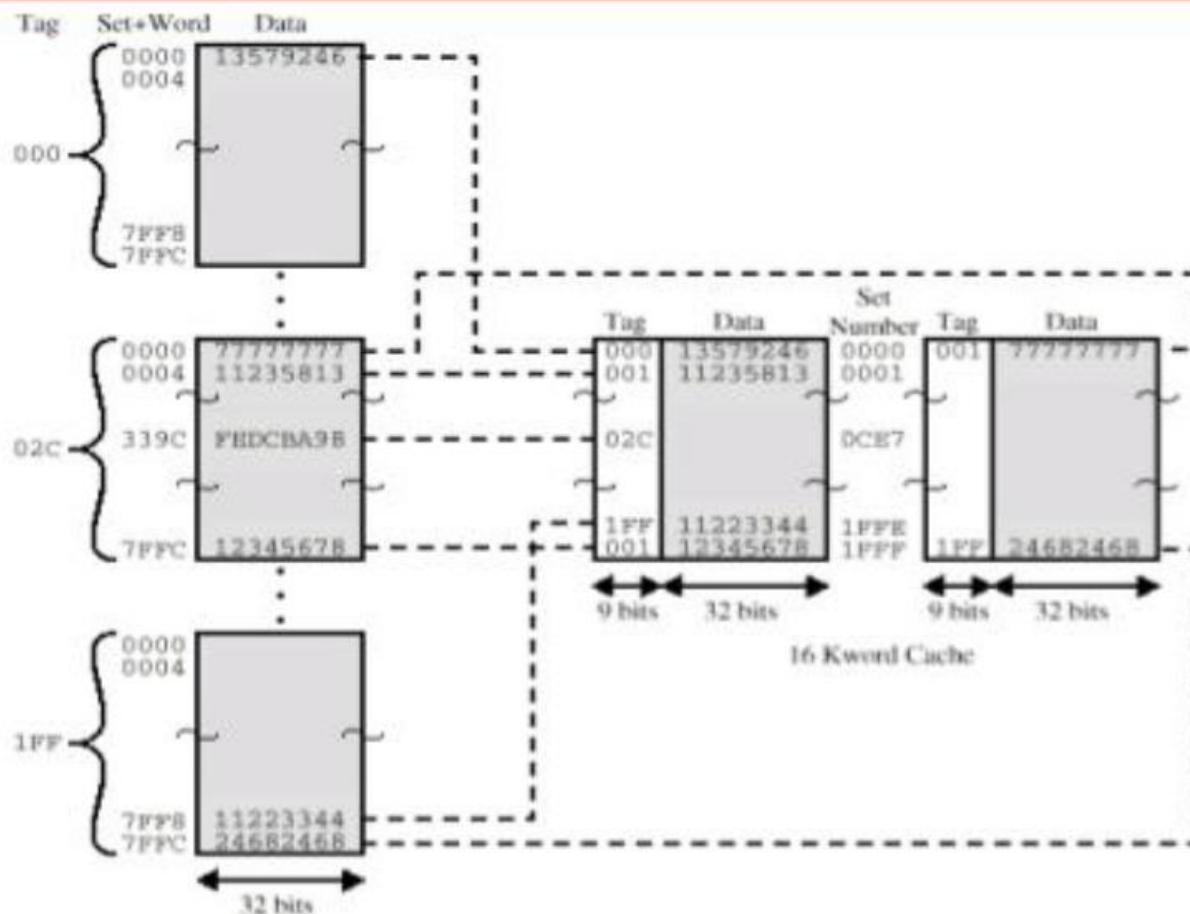
12(16-4)

4 (2^4)

Working of Fully Associative Mapping

- "Tag" field identifies one of the $2^{12} = 4096$ memory lines;
- All the cache tags are searched to find out whether or not the Tag field matches one of the cache tags.
- If so, we have a hit, and if not there's a miss and we need to replace one of the cache lines by this line before reading or writing into the cache.
- (The "Word" field again selects one from among 16 addressable words (bytes) within the line.)

Two Way Set Associative Mapping Example



SET ASSOCIATIVE MAPPING(2-way)

- Cache- 128 blocks of 16 words each
- $128 * 16 = 2048$ **approx 2 KB**
- Set divided into 2 ($128 / 2$) =64
- Main Memory – 4K blocks of 16 words each
- $4K * 16 = 64000$ **approx 64 KB**

Tag	Set	Word
-----	-----	------

6($16-(4+6)$)

6(2^6)

4 (2^4)

Working

- "Tag" field identifies one of the $2^6 = 64$ different memory lines in each of the $2^6 = 64$ different "Set" values.
- Since each cache set has room for only two lines at a time, the search for a match is limited to those two lines (rather than the entire cache).
- If there's a match, we have a hit and the read or write can proceed immediately.
- Otherwise, there's a miss and we need to replace one of the two cache lines by this line before reading or writing into the cache. (The "Word" field again select one from among 16 addressable words inside the line.)
- In set-associative mapping, when the number of lines per set is n , the mapping is called n -way associative. For instance, the above example is 2-way associative.

Direct Mapping

Cache Line Table

- Cache line Main Memory blocks held
0, m, 2m, 3m... $2s-m$
- 0 $1, m+1, 2m+1 \dots 2s-m+1$
- 1 $m-1, 2m-1, 3m-1 \dots 2s-1$
- $m-1$

K J Somaiya College of Engineering

Direct mapping

- The direct mapping technique is simple and inexpensive to implement.
- Its main disadvantage is that there is a fixed cache location for any given block.
- Thus, if a program happens to reference words repeatedly from two different blocks that map into the same line, then the blocks will be continually swapped in the cache, and the hit ratio will be low (a phenomenon known as *thrashing*).

Direct Mapping pros & cons

- Simple
- Inexpensive
- Fixed location for given block
 - If a program accesses 2 blocks that map to the same line repeatedly, cache misses are very high

Associative Mapping

- A main memory block can load into any line of cache
- Memory address is interpreted as tag and word
- Tag uniquely identifies block of memory
- Every line's tag is examined for a match
- Cache searching gets expensive

Set Associative Mapping

- Cache is divided into a number of sets
- Each set contains a number of lines
- A given block maps to any line in a given set
 - e.g. Block B can be in any line of set i
- e.g. 2 lines per set
 - 2 way associative mapping
 - A given block can be in one of 2 lines in only one set

Problem statement

Consider a cache consisting of 256 blocks of 16 words each for a total of 4096(4KB) words and assume that the main memory is addressable by a 16 bit address and it consists of 4KB blocks of 16 words.

How many bits are there in each of the TAG,BLOCK/SET and WORD field for Direct Mapping , Fully Associative and 2-way set associative techniques?

Problem statement

A block set associative cache memory consists of 128 blocks divided into four block sets. The main memory consists of 16,384 blocks and each block contains 256 words.

- i) How many bits are required for addressing the main memory?
- ii) How many bits are needed to represent the TAG, SET and WORD fields?

Problem statement

A block set associative cache memory consists of 64 blocks divided into four block sets. The main memory consists of 4096 blocks and each block contains 128 words.

- i) How many bits are there in main memory?
- ii) How many bits are needed to represent the TAG, SET and WORD fields?

Replacement Algorithms

- Once the cache has been filled, when a new block is brought into the cache, one of the existing blocks must be replaced.
- -Four of the most common-
- Least recently used (LRU)
- First-in-first-out (FIFO)
- Least frequently used (LFU)

K J Somaiya College of Engineering
Write Policy

- When a block that is resident in the cache is to be replaced, there are two cases to consider.
 - Old block in the cache has not been altered.
 - overwritten without first writing out the old block-
 - One write operation has been performed.
 - main memory must be updated-

Write Policy

- A variety of write policies, is possible.
- -Simplest technique called **write through**-
- All write operations are made to main memory & cache, ensuring Main Memory is always valid-

-Disadvantage-

- Generates substantial memory traffic.

Alternative method -write back

-Minimizes memory writes.

-Updates are made only in the cache.

When an update occurs, a **dirty bit**, or **use bit**, associated with the line is set .Then, when a block is replaced, it is written back to main memory only if dirty bit is set.

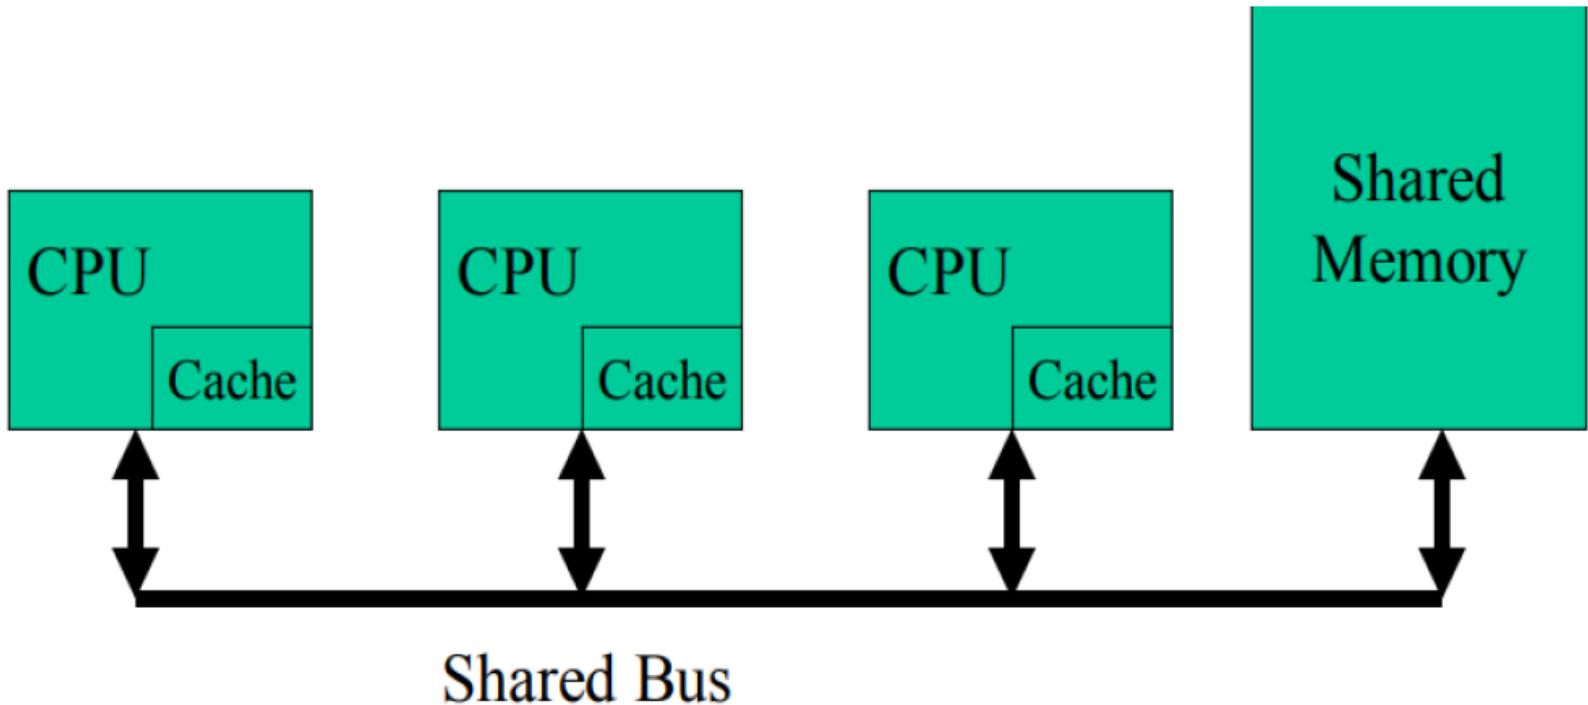
Write through

- All writes go to main memory as well as cache
- Multiple CPUs can monitor main memory traffic to keep local (to CPU) cache up to date
- Lots of traffic
- Slows down writes

Write back

- **Updates initially made in cache only**
- Update bit for cache slot is set when update occurs
- If block is to be replaced, write to main memory only if **update bit is set**
- Other caches get out of sync
- I/O must access main memory through cache

CACHE COHERENCE



Cache Coherence

- Problem - multiple copies of same data in different caches
- Can result in an **inconsistent** view of memory
 - Write through
 - Write back policy
 - Write invalidate
 - Write Update

Write policies & cache coherence Approaches

- **Write back:** Write operations are usually made only to the cache. Main memory is only updated when the corresponding cache line is flushed from the cache.
- **Write through:** All write operations are made to main memory as well as to the cache, ensuring that main memory is always valid.
- **Cache coherence approaches** -

- Divided into software and hardware approaches.
- Some adopt a strategy involving both software and hardware elements

Software Solutions

- **Compiler and operating system** deal with problem
- Overhead transferred to compile time
- Compiler **marks** data likely to be changed and OS prevents such data from being cached
- Design complexity transferred from hardware to software

Hardware Solution

- Cache coherence protocols
- Dynamic recognition of potential problems
- Run time
- More **efficient** use of cache
- **Transparent** to programmer
- **Snoopy protocols**(to maintain cache consistency)

Hardware solutions

- Hardware-based solutions are generally referred to as cache coherence protocols.
- Because the problem is only dealt with when it actually arises, there is more effective use of caches, leading to improved performance over a software approach.
- These approaches are transparent to the programmer and the compiler, reducing the software development burden.
- Hardware schemes differ in a number of particulars, including where the state information about data lines is held, how that information is organized, where coherence is enforced, and the enforcement mechanisms. In general, hardware schemes can be divided into two categories: directory protocols and snoopy protocols.

Snoopy Protocols

- Distribute cache coherence responsibility among **cache controllers**
- Cache recognizes that a line is shared
- Updates announced to other caches
- Suited to bus based multiprocessor
- Increases bus traffic

K J Somaiya College of Engineering
Snoopy Protocol

- **2 Basic Approaches of SNOOPY PROTOCOL**

-1. Write Invalidate

-2. Write update

Write Invalidate

- **Multiple readers, one writer**
- When a write is required, all other caches of the line are invalidated
- Writing processor then has **exclusive access** until line required by another processor
- State of every line is marked as modified, exclusive, shared or invalid
- MESI protocol

Write Update

- Multiple readers and writers
- Updated word is distributed to all other processors
- Some systems use an adaptive mixture of both solutions

MESI Protocol

Commonly implemented for Cache coherence

MESI protocol -four states that a cache line may be in:

- **Modified**
- **Exclusive**
- **Shared**
- **Invalid**

MESI protocol

- **I**nvalid: This cache line is not valid
- **E**xclusive: This **cache** has the only copy of the data. The **memory** is valid.
- **S**hared: More than one cache is holding a copy of this line. The memory copy is valid.
- **M**odified: The line has been modified. The **memory** copy is invalid.

Interleaved And Associative Memory

Interleaved Memory

- **Interleaved memory** is a design made to compensate for the relatively slow speed of DRAM
- **Spreads memory** addresses evenly across
- Contiguous memory reads and writes
- Resulting in **higher memory throughputs** due to reduced waiting.

Chip 00

C[0]	M[0]
C[1]	M[4]
C[1023]	M[4092]

Chip 01

C[0]	M[1]
C[1]	M[5]
C[1023]	M[4093]

Chip 10

C[0]	M[2]
C[1]	M[6]
C[1023]	M[4094]

Chip 11

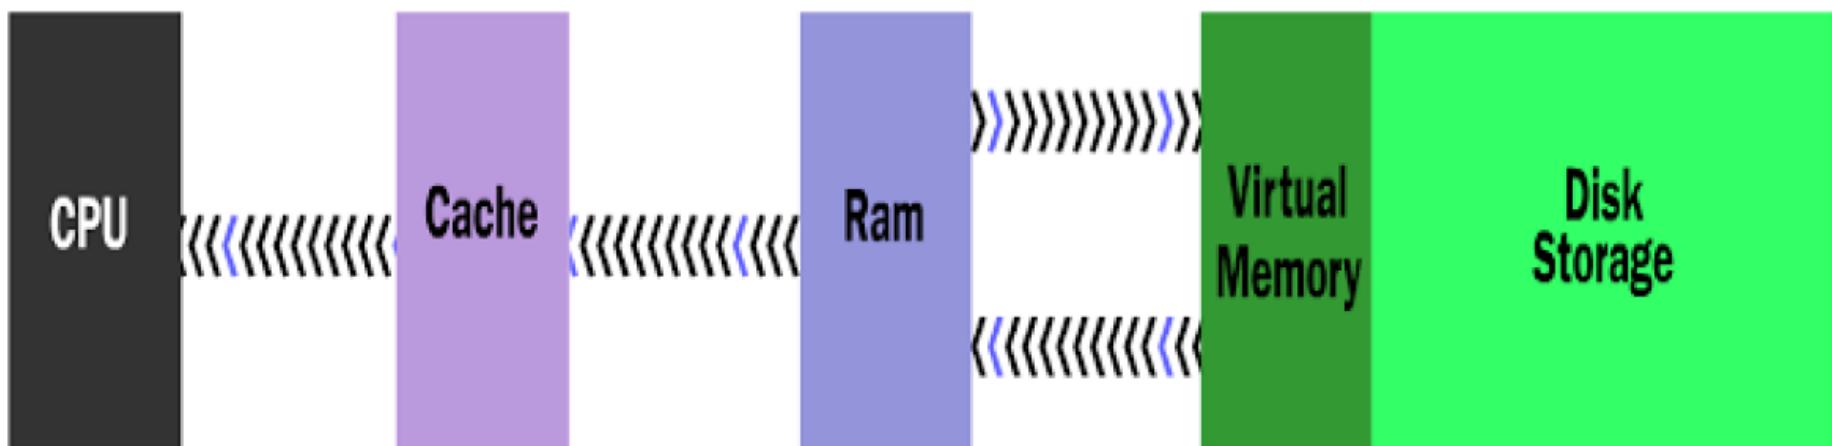
C[0]	M[3]
C[1]	M[7]
C[1023]	M[4095]

Associative Memory

- Content-addressed or associative memory-memory is **accessed by its content** (as opposed to an explicit address).
- **Reference clues** are "associated" with actual memory contents until a desirable match (or set of matches) is found.
- Humans retrieve information best when it can be linked to other related information.

Virtual Memory

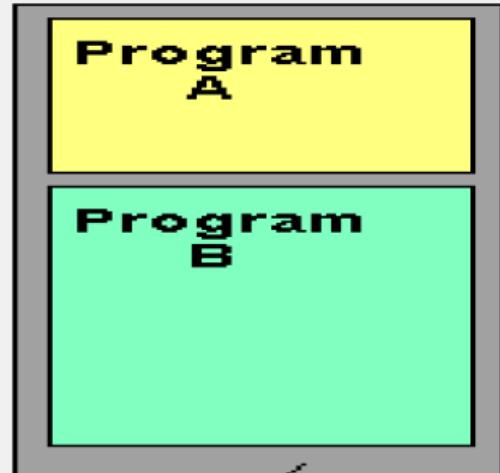
Memory Management



© 2000 How Stuff Works, Inc

NO VIRTUAL MEMORY

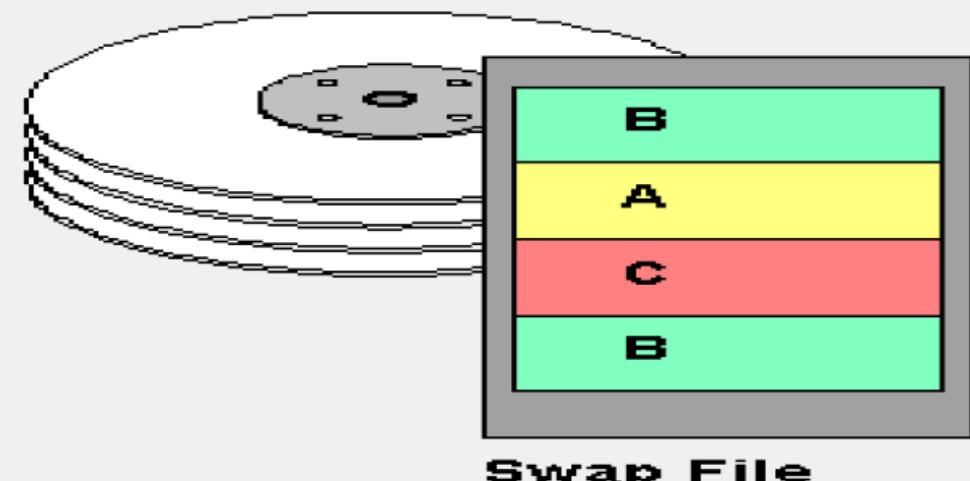
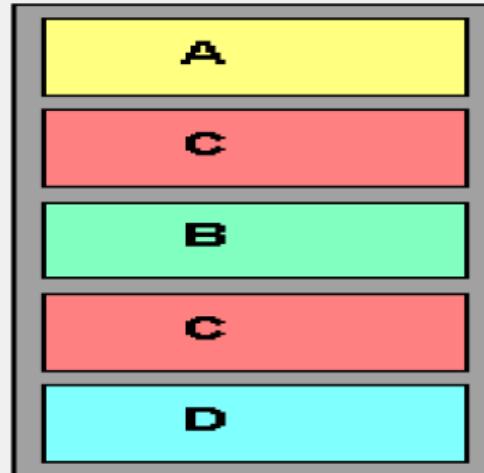
Real Memory



No more
programs fit.

VIRTUAL MEMORY COMPUTER

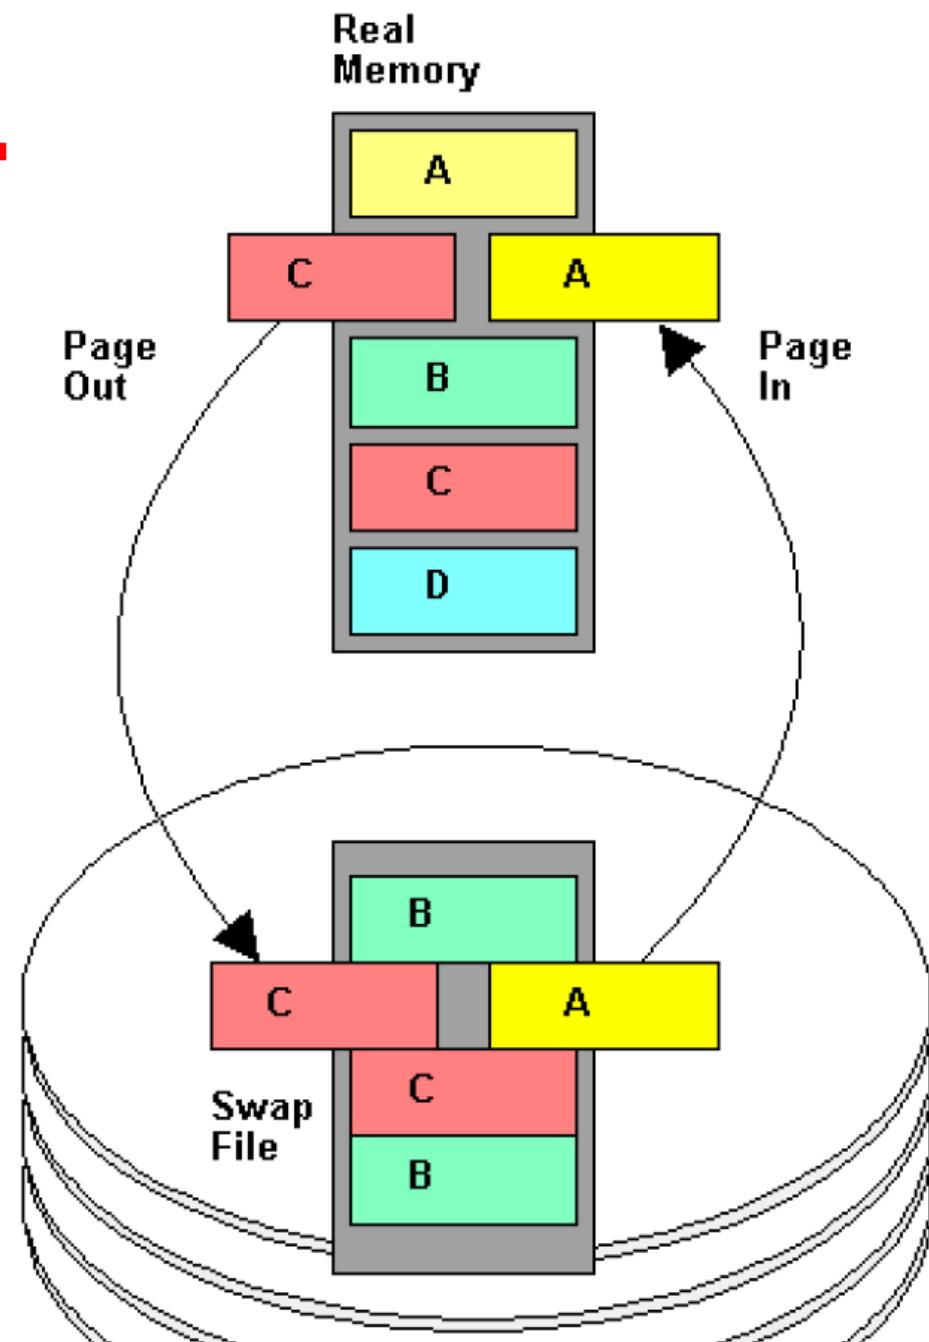
Real Memory



Swap File

• Virtual memory

Allows more programs to be opened simultaneously by using the hard disk as temporary storage of memory pages.



VIRTUAL MEMORY

- 32 or 64MB of RAM available for CPU usage.
- Users expect all their programs to run at once.
- Ex Email program,a Web browser and word processor(all in RAM simultaneously)
- Find RAM for areas that have not been used recently and copy them onto the hard disk

VIRTUAL MEMORY

- Frees up space in RAM to load the new application.
- **Copying** happens automatically(feels like unlimited RAM space)
- Hard disk space is much cheaper than RAM chips, thus has a economic benefit.
- **Read/write speed of a hard drive & technology** is not geared toward accessing small pieces of data at a time.

VIRTUAL MEMORY

- **Operating system** has to constantly swap information back and forth between RAM and the hard disk.
- **Thrashing**- computer feels incredibly slow.

PAGING

- Unequal fixed size /Variable Size partitions(Inefficient)
- Primary memory is divided into **small equal fixed sized partitions** (256, 512, 1K) called **page frames**.
- Process are divided into **same sized blocks(pages)** called **paging**.
- Recently referenced pages in the memory.
- Need a **page table** to this management.

Frame number	Main memory
0	
1	
2	
3	
4	
5	
6	
7	
8	
9	
10	
11	
12	
13	
14	

(a) Fifteen Available Pages

Frame number	Main memory
0	A.0
1	A.1
2	A.2
3	A.3
4	
5	
6	
7	
8	
9	
10	
11	
12	
13	
14	

(b) Load Process A

Frame number	Main memory
0	A.0
1	A.1
2	A.2
3	A.3
4	B.0
5	B.1
6	B.2
7	
8	
9	
10	
11	
12	
13	
14	

(b) Load Process B

Frame number	Main memory
0	A.0
1	A.1
2	A.2
3	A.3
4	B.0
5	B.1
6	B.2
7	C.0
8	C.1
9	C.2
10	C.3
11	
12	
13	
14	

(d) Load Process C

Frame number	Main memory
0	A.0
1	A.1
2	A.2
3	A.3
4	
5	
6	
7	C.0
8	C.1
9	C.2
10	C.3
11	
12	
13	
14	

(e) Swap out B

Frame number	Main memory
0	A.0
1	A.1
2	A.2
3	A.3
4	D.0
5	D.1
6	D.2
7	C.0
8	C.1
9	C.2
10	C.3
11	D.3
12	D.4
13	
14	

(f) Load Process D

Figure 7.9 Assignment of Process Pages to Free Frames

Page Table Sample

Figure 5-8. Format of a Linear Address

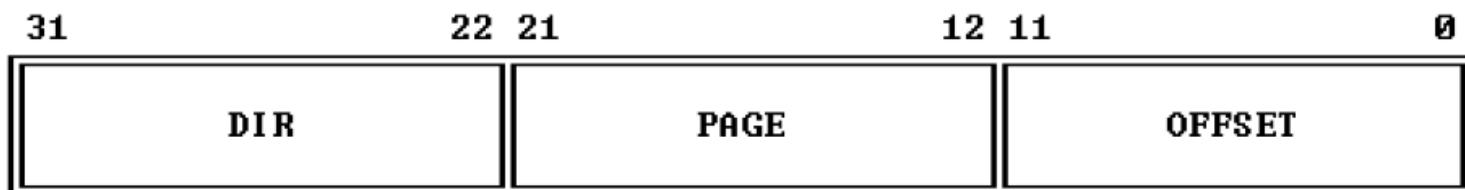
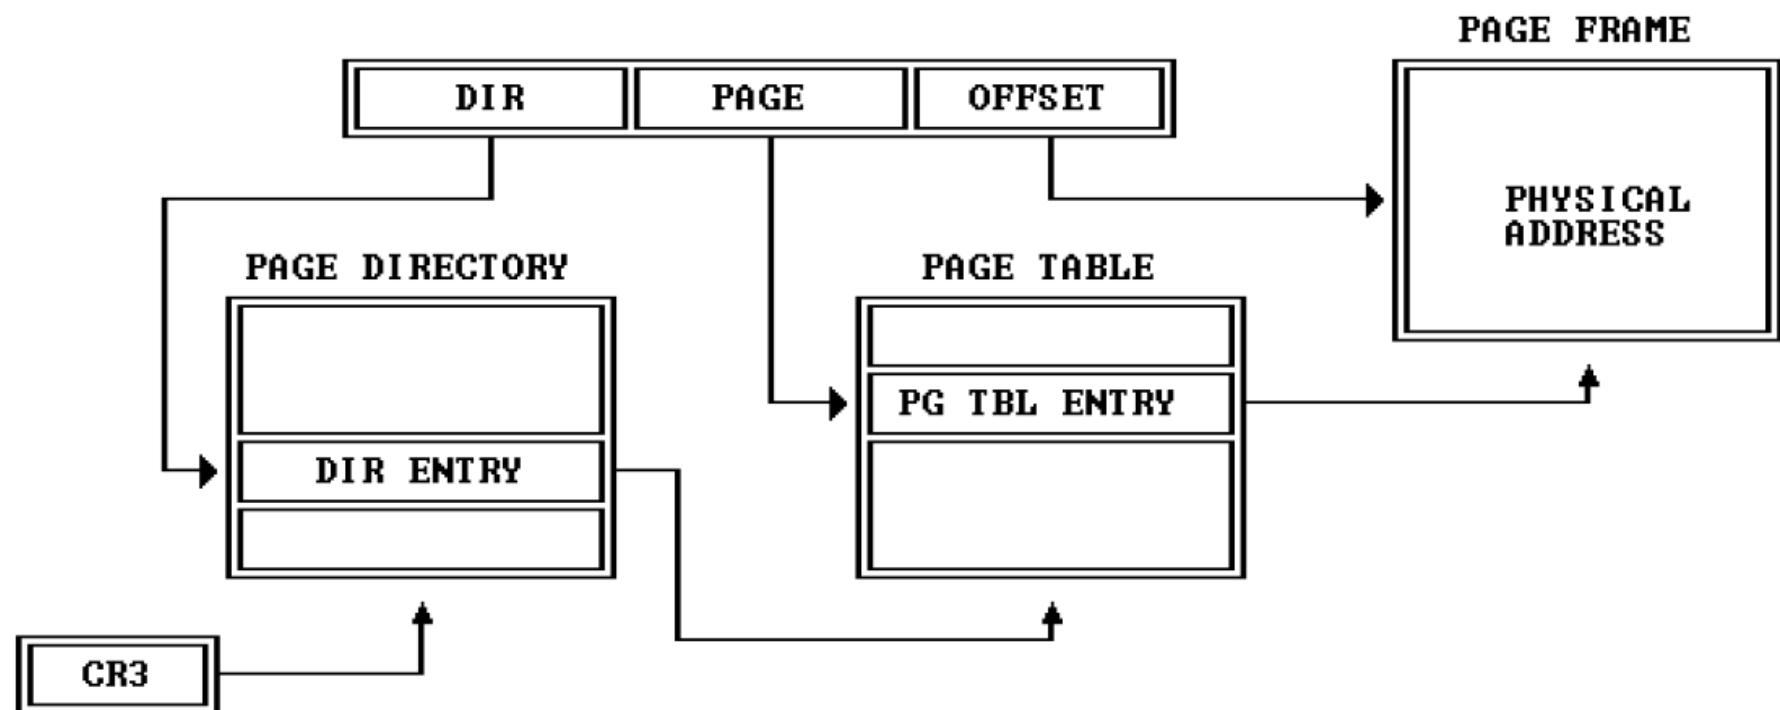


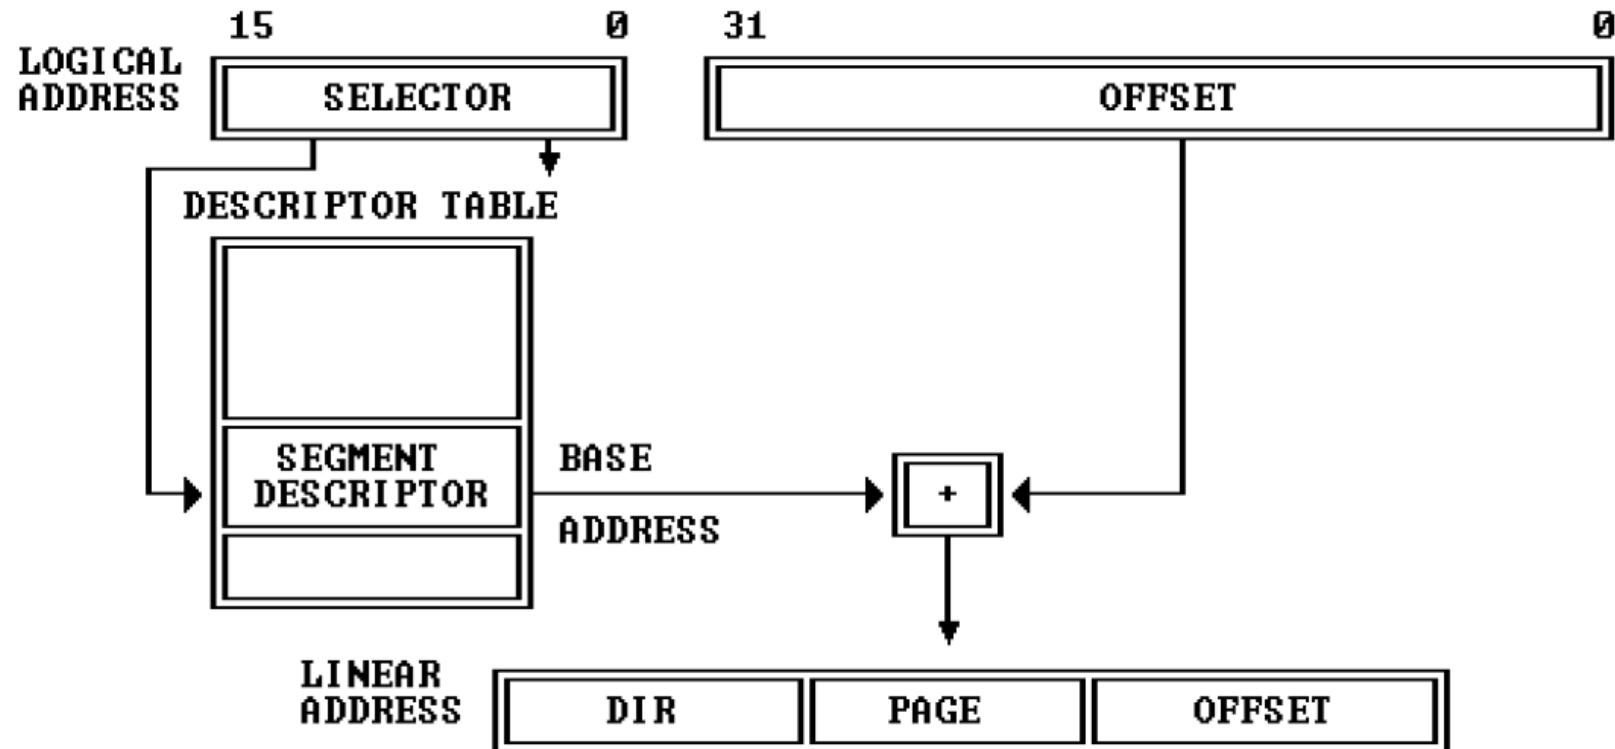
Figure 5-9. Page Translation



SEGMENTATION

- Paging → **internal fragmentation**.
- Segmentation maps segments representing data structures, modules, etc. into **variable partitions**.
- Nor contiguous memory blocks neither all segments of a process are loaded at a time.
- We need a **segment table** very much like a page table.

Figure 5–2. Segment Translation



Main Memory Allocation

- Memory is divided into set of contiguous locations called regions/segments/pages
- Store blocks of data
- Placement of blocks of information in memory is called **Memory Allocation**
- **Memory Management Systems** keeps information in a table containing available and free slots

Replacement Algorithms

- Hardware implemented algorithm (speed)
- Least Recently used (LRU)
 - Pick the slot that hasn't been used in the longest time.
- First in first out (FIFO)
 - replace block that has come into cache first
- Random
- OPT-Optimal(Future)

FIFO, LRU, OPT

1) 1 , 6 , 4 , 5 , 1 , 4, 3, 2, 1, 2, 1, 4, 6, 7, 4

FIFO → 7 - 4 - 6

LRU → 4 - 6 - 7

OPT → 7 - 6 - 4 (Conflict resolved using LRU)

2) 2 , 3 , 2 , 1 , 5 , 2 , 4 , 5 , 3 , 2 , 5 , 2

FIFO → 3 - 2 - 5

LRU → 3 - 5 - 2

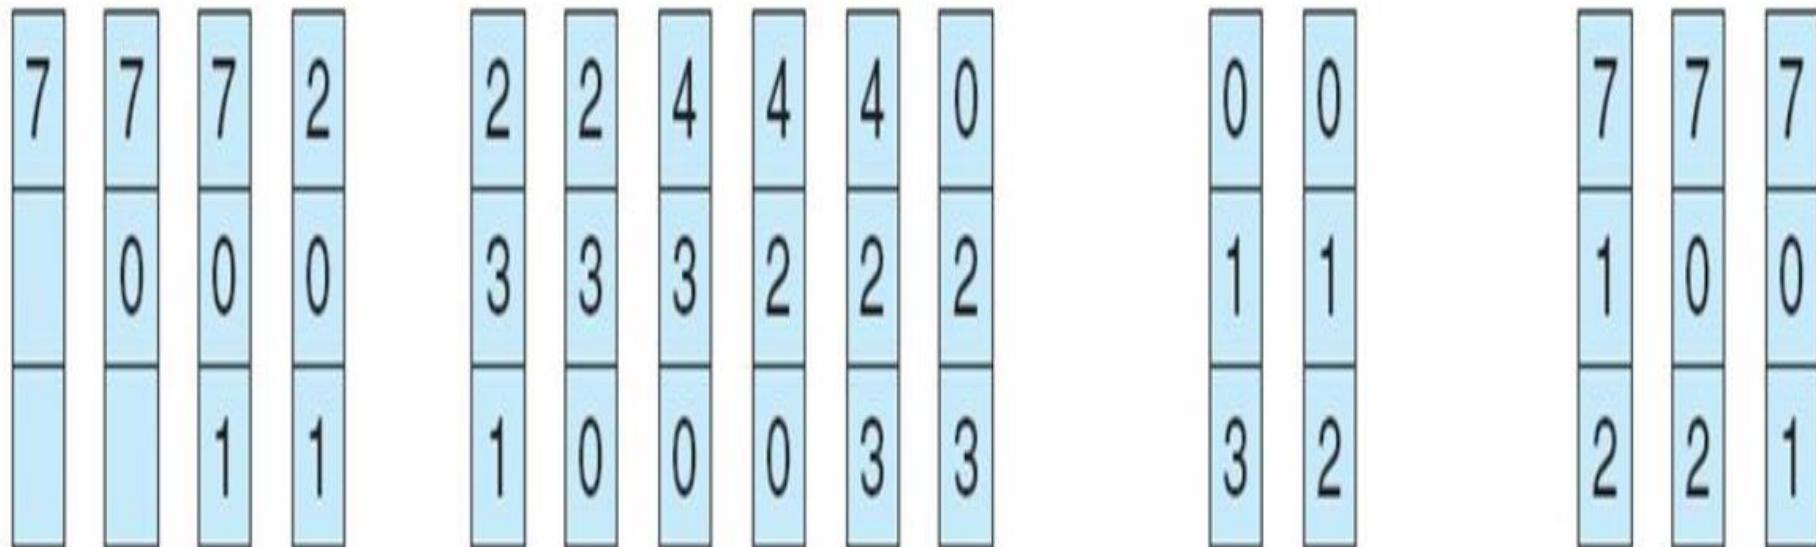
OPT → 2 - 3 - 5

7,0,1,2,0,3,0,4,2,3,0,3,2,1,2,0,1,7,0,1

FIFO Page Replacement

reference string

7 0 1 2 0 3 0 4 2 3 0 3 2 1 2 0 1 7 0 1

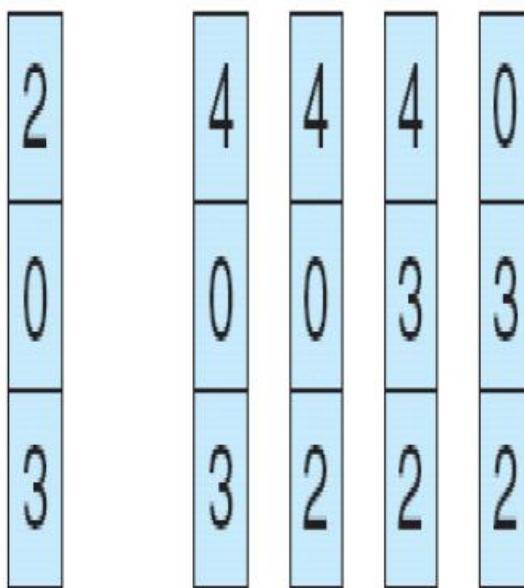
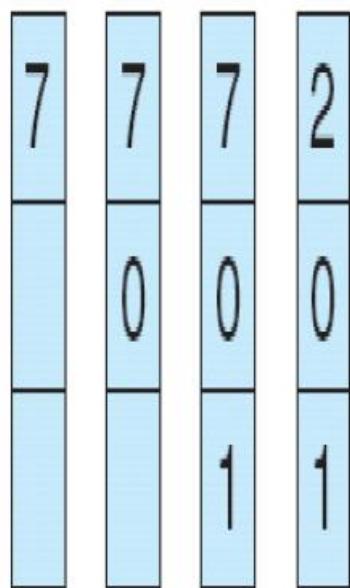


page frames

LRU Page Replacement

reference string

7 0 1 2 0 3 0 4 2 3 0 3 2 1 2 0 1 7 0 1

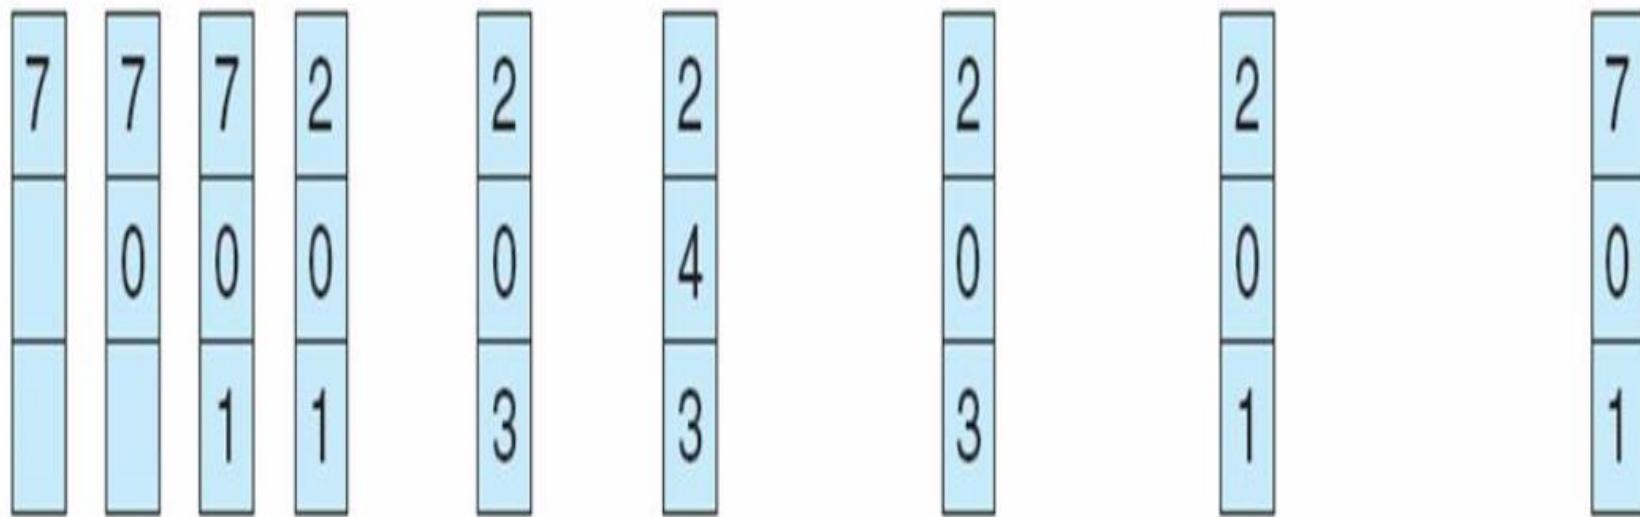


page frames

Optimal Page Replacement

reference string

7 0 1 2 0 3 0 4 2 3 0 3 2 1 2 0 1 7 0



page frames

Secondary Storage

- Magnetic disks
- Floppy disks
- Magnetic Tape
- RAID
- Optical Memory
- CD-ROM
- DVD

RAID

(Redundant Array Of Independent Disks)

Secondary Storage-Very Slow-

-Additional gain in performance is obtained by using multiple parallel components.

In disk storage - arrays of disks are developed that operate independently and in parallel.

With Multiple disks,-separate I/O requests handled in parallel, as long as the data required is on separate disks.

RAID

- With use of multiple disks- wide variety of ways in which the data organized and redundancy added to improve reliability.
- So standardized scheme for multiple-disk database design, known as **RAID –Redundant Array of Independent Disks**.

RAID scheme consists of **seven levels**, zero through six.

RAID

These levels designate different design architectures that share three common characteristics:

- **1.** RAID is a set of physical disk drives viewed by the operating system as a single logical drive.
- **2.** Data are distributed across the physical drives of an array in a scheme known as striping.
- **3.** Redundant disk capacity is used to store parity information, which guarantees data recoverability in case of a disk failure.

RAID Levels 0 - 6

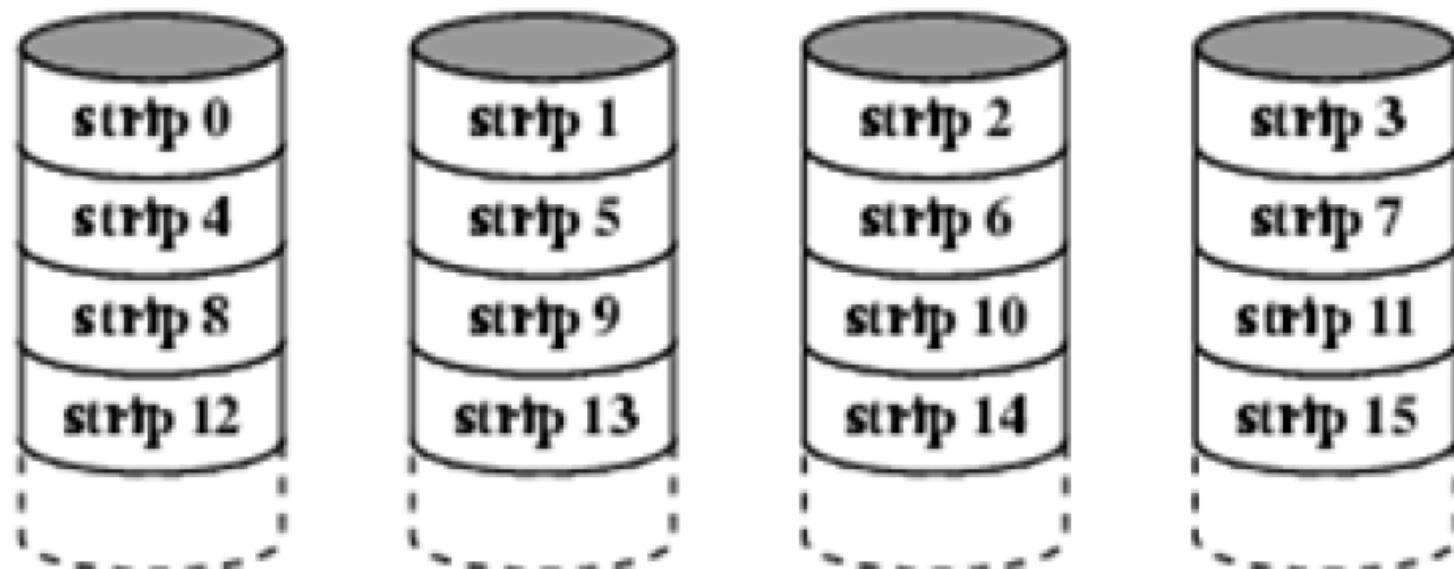
REDUNDANT ARRAY OF INDEPENDENT DISKS

- Storage is an important consideration when setting up a server.
- Almost all of the important information that you and your users care about will at one point be written to a storage device to save for later retrieval.
- Single disks can serve you well if your needs are straight forward.
- However, if you have more complex redundancy or performance requirements, solutions like RAID can be helpful.

RAID

- The RAID strategy employs –
- Multiple disk drives and
- Distributes data in these disk drives to enable simultaneous access to data from multiple drives.
- Improving I/O performance and allowing increases in capacity.

RAID Level 0- Non Redundant



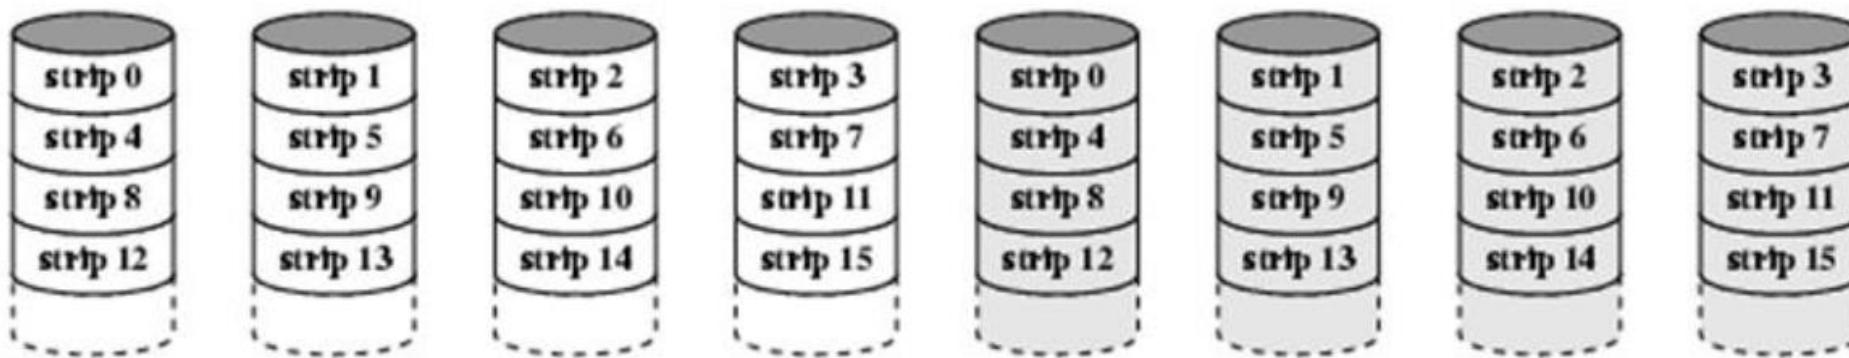
(a) RAID 0 (non-redundant)

RAID LEVEL-0

(*FOR HIGH DATA TRANSFER CAPACITY*)

- NO Redundancy
- Improved Performance & Capacity .
- Low cost.
- User & system data distributed across all disks.
- Parallel issue of request-Reducing I/O Queuing time.
- The disk s divided into strips (like blocks or frames in MM).
- Best performance-but no fault tolerance.

RAID Level -1 Mirrored

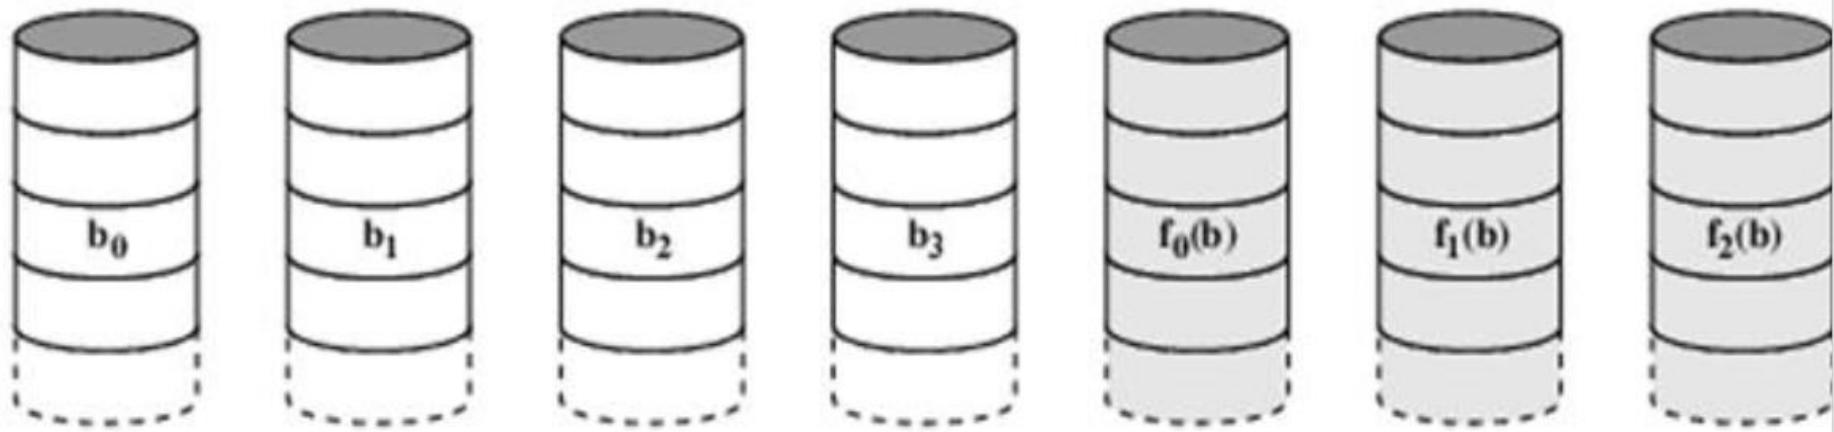


(b) RAID 1 (mirrored)

RAID LEVEL-1

- Redundancy of data.
- Redundancy achieved by duplicating data.
- Every disk in the array has a mirror disk with same data.
- A read request can be serviced by either of the 2 disks-Improving read performance.
- Write request is done in parallel.
- Failure recovery very easy-data still present in second drive.
- Costly-twice disk space.

RAID Level 2- Hamming Code

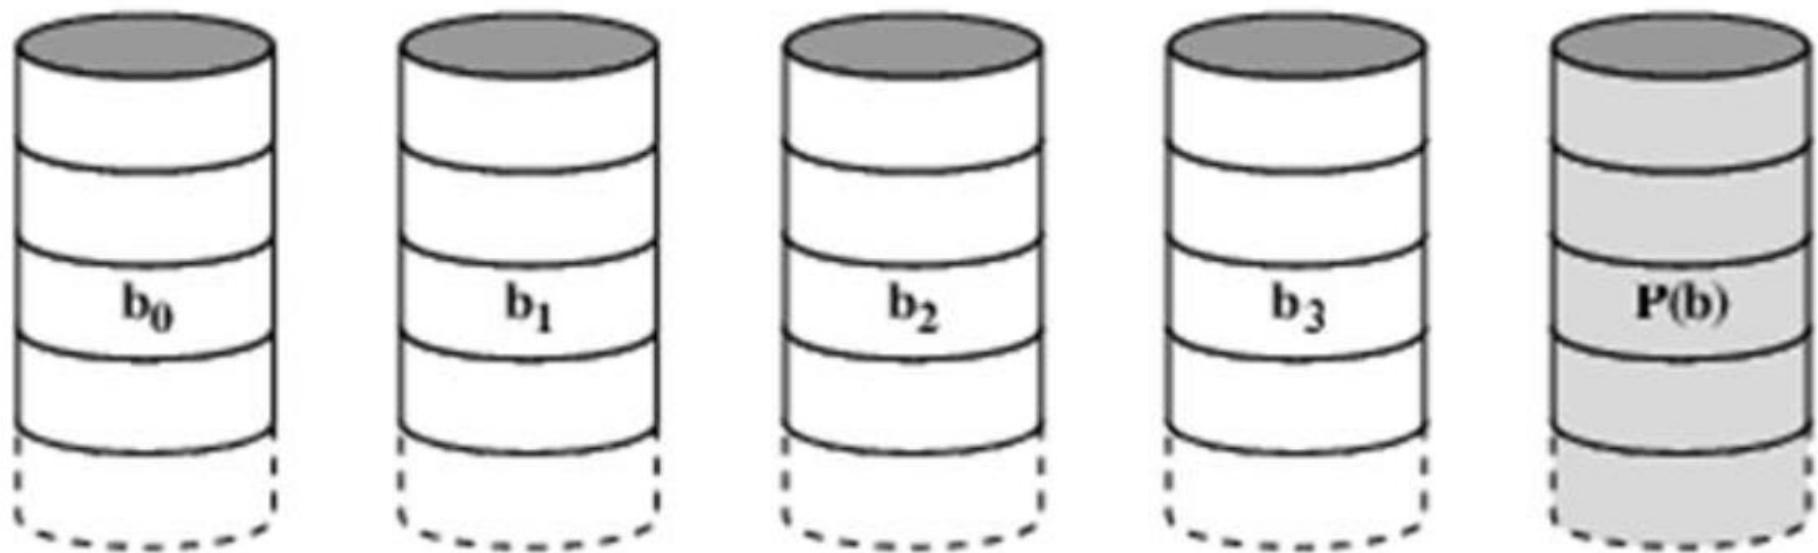


(c) RAID 2 (redundancy through Hamming code)

RAID LEVEL-2

- Parallel access technique-All member disk takes part in execution of every I/O request.
- Data striping is used-as in other levels-with very smaller strips.
- Error correcting code –calculated across corresponding bits on each data disk and stored in multiple parity disks.
- Some discs store error correcting & parity info.
- Hamming code used-error correction(correct single bit errors & detect double bit errors).
- Fewer disks than RAID 1-but still costly
- No: of redundant disks proportional to $\log(\text{no : of data disk})$.

RAID Level 3 –Bit Interleaved Parity



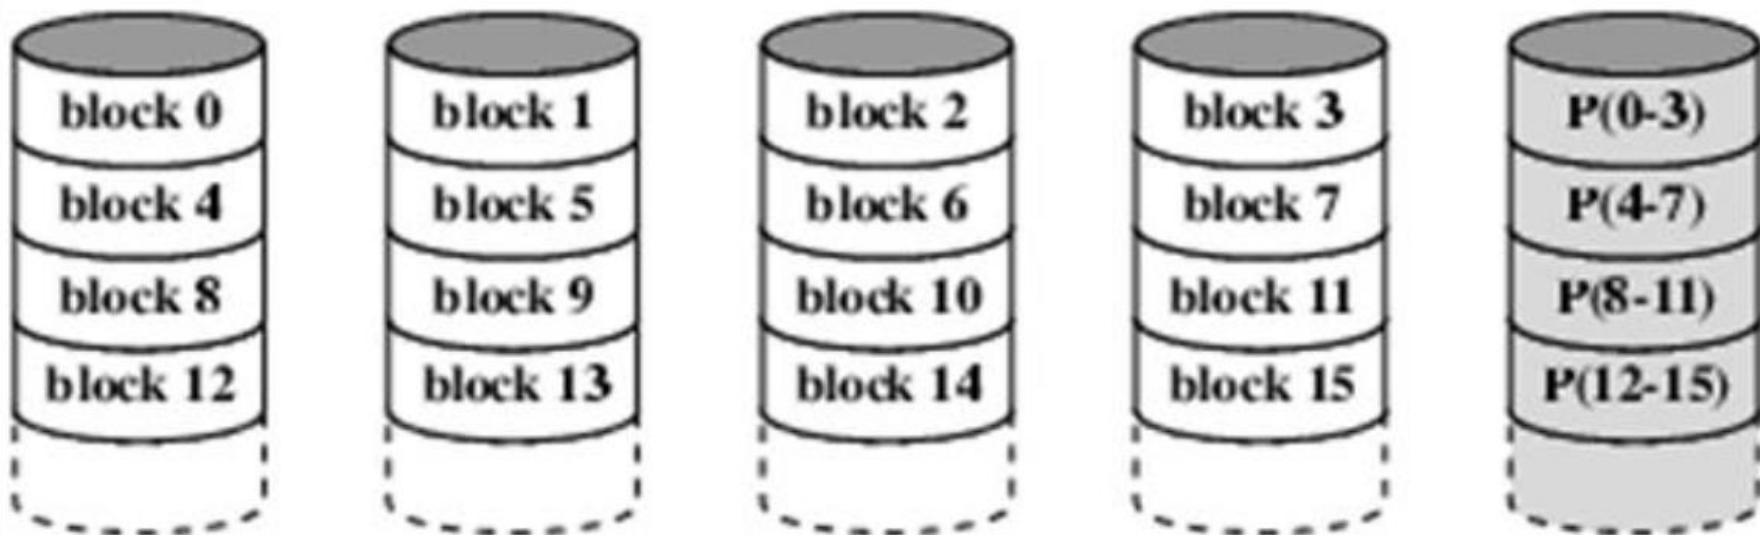
(d) RAID 3 (bit-interleaved parity)

RAID-3

- Organized in a similar fashion to RAID 2-
- Difference –This requires only a single disk to storing parity, no matter how large the disk array.
- Employs parallel access, with data distributed in small strips.
- A simple parity bit is computed for the set of individual bits in the same position on all of the data disks.
- ***REDUNDANCY***- In the event of a drive failure, the parity drive is accessed and data is reconstructed from the remaining disks in the array.
- Once the failed drive is replaced, the missing data can be restored on the new drive and operation resumed.

- **PERFORMANCE** -Data are striped in very small strips, RAID 3 can achieve very high data transfer rates.
- Any I/O request will involve the parallel transfer of data from all of the data disks.
- Only one I/O request can be executed at a time.
- Thus, in a transaction-oriented environment, performance suffers.

RAID Level 4 - Block level parity



(e) RAID 4 (block-level parity)

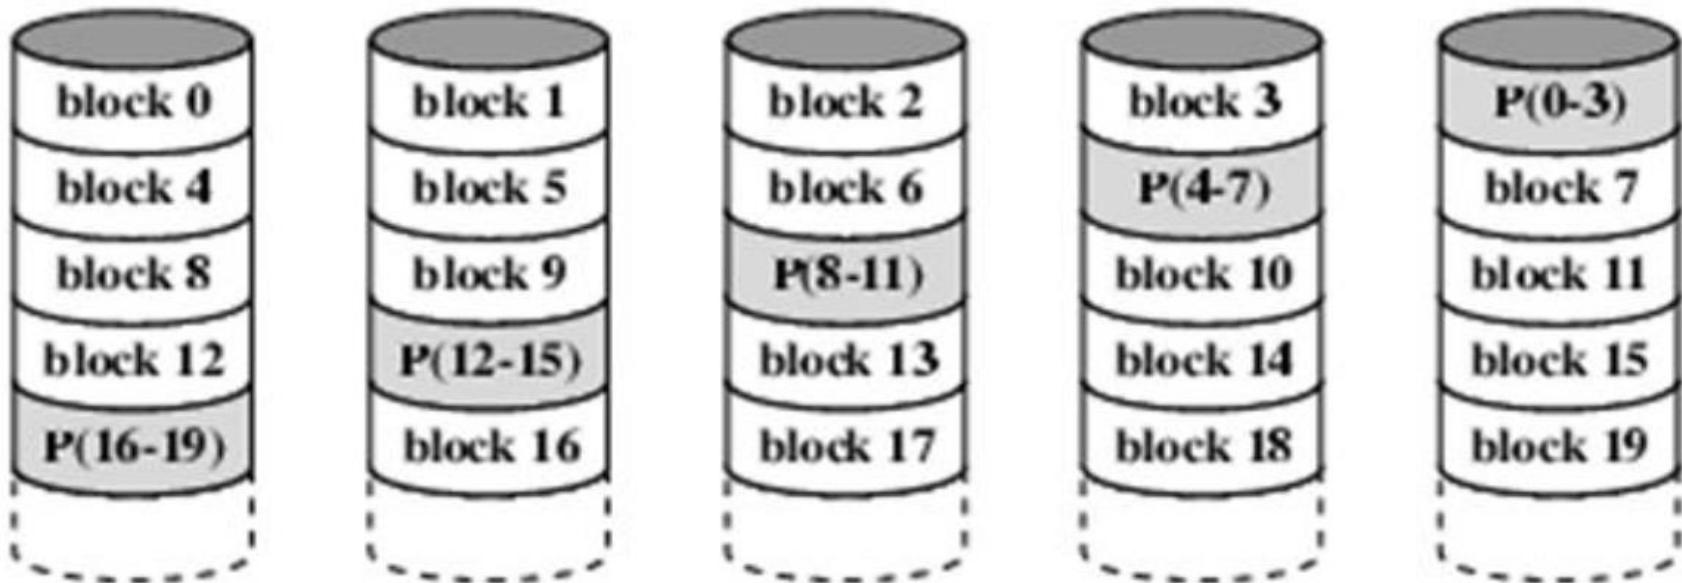
RAID_4

- Separate I/O requests can be satisfied in parallel.
- Uses large stripes-means a user can read records from any single drive.
- Suitable for applications- that require high I/O request rates and less suited for applications -that require high data transfer rates.
- A bit-by-bit parity strip is calculated across corresponding strips on each data disk, and the parity bits are stored in the corresponding strip on the parity disk.

RAID-4

- Each update requires, -updating the user data and also the parity drive.
- To calculate the new parity, -read the old user strip and the old parity strip.
- Then update these two strips with the new data and the newly calculated parity.
- Thus, each strip write involves two reads and two writes.

RAID Level 5- Block level Distributed Parity

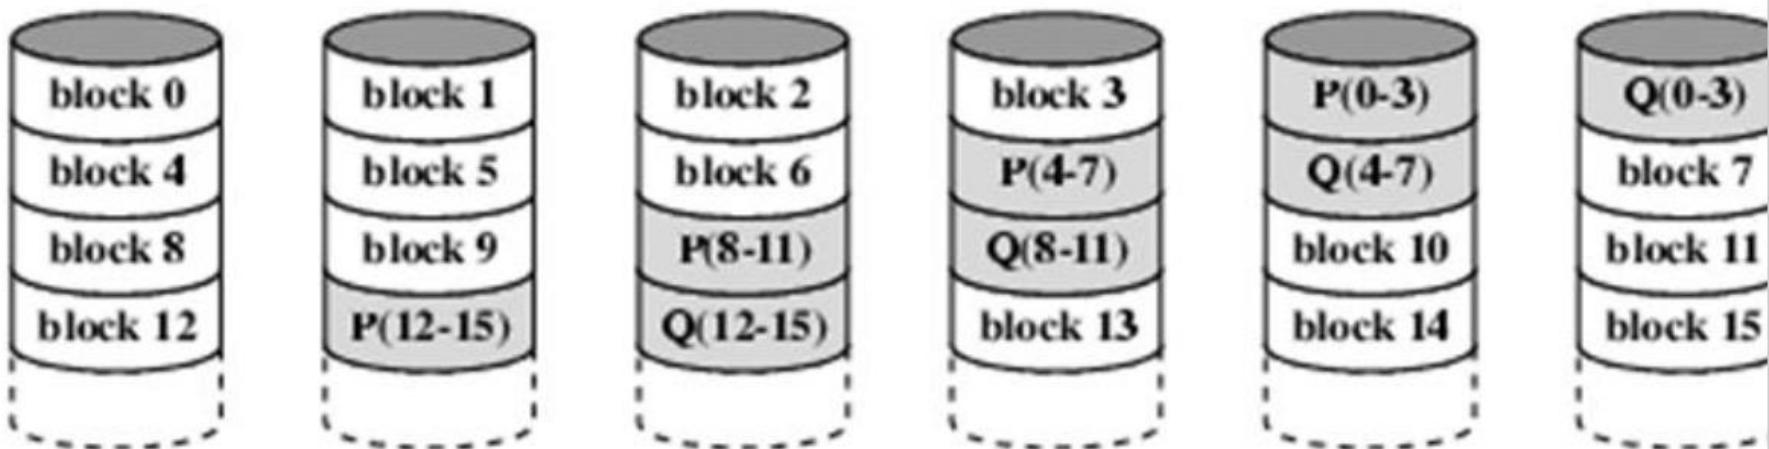


(f) RAID 5 (block-level distributed parity)

RAID-5

- RAID 5 is organized in a similar fashion to RAID 4.
- The difference - parity information is striped across each drive, enabling the array to function, even if one drive were to fail.
- A typical allocation is a round-robin scheme.
- For an n-disk array, the parity strip is on a different disk for the first n stripes, and the pattern then repeats.

RAID Level 6- Dual Redundancy



(g) RAID 6 (dual redundancy)

RAID-6

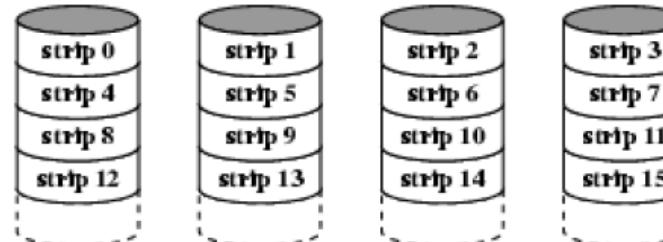
Here, two different parity calculations are carried out and stored in separate blocks on different disks.

- 1) Exclusive-OR calculation used in RAID 4 and 5.
 - 2) An independent data check algorithm.
- In RAID-6, if user data require N disks consists of N +2 disks.
 - Data check algorithm makes it possible to regenerate data even if two disks containing user data fail.
 - The advantage of RAID 6 is that it provides extremely high data availability.
 - RAID 6 incurs a substantial write penalty, because each write affects two parity blocks.
 - RAID 6 controller suffer more than a 30% drop in overall write performance compared with a RAID 5 implementation.

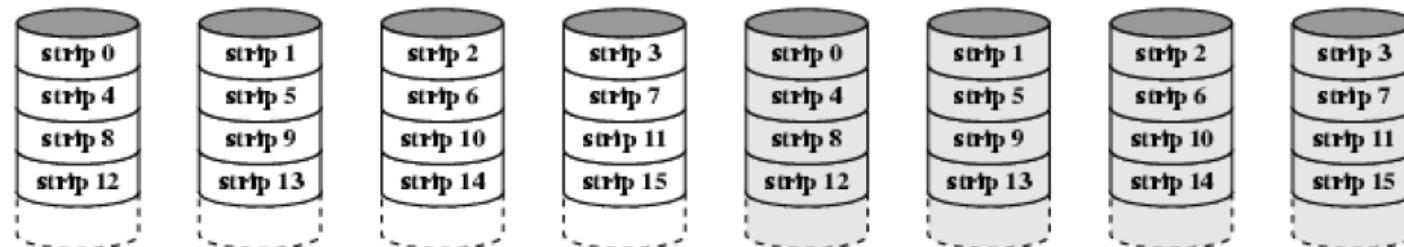
Category	Level	Description	Disk Required	Data Availability	Large I/O Data Transfer Capacity	Small I/O Request Rate
Striping	0	Nonredundant	N	Lower than single disk	Very high	Very high for both read and write
Mirroring	1	Mirrored	$2N$	Higher than RAID 2, 3, 4, or 5; lower than RAID 6	Higher than single disk for read; similar to single disk for write	Up to twice that of a single disk for read; similar to single disk for write
Parallel access	2	Redundant via Hamming code	$N + m$	Much higher than single disk; comparable to RAID 3, 4, or 5	Highest of all listed alternatives	Approximately twice that of a single disk
	3	Bit-interleaved parity	$N + 1$	Much higher than single disk; comparable to RAID 2, 4, or 5	Highest of all listed alternatives	Approximately twice that of a single disk
Independent access	4	Block-interleaved parity	$N + 1$	Much higher than single disk; comparable to RAID 2, 3, or 5	Similar to RAID 0 for read; significantly lower than single disk for write	Similar to RAID 0 for read; significantly lower than single disk for write
	5	Block-interleaved distributed parity	$N + 1$	Much higher than single disk; comparable to RAID 2, 3, or 4	Similar to RAID 0 for read; lower than single disk for write	Similar to RAID 0 for read; generally lower than single disk for write
	6	Block-interleaved dual distributed parity	$N + 2$	Highest of all listed alternatives	Similar to RAID 0 for read; lower than RAID 5 for write	Similar to RAID 0 for read; significantly lower than RAID 5 for write

N = number of data disks; m proportional to $\log N$

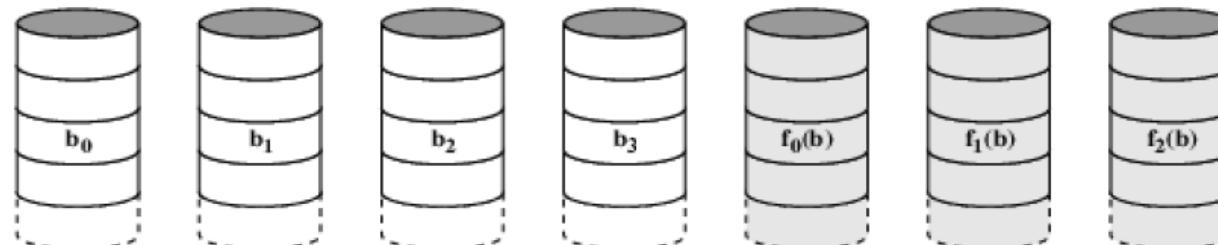
RAID 0, 1, 2 – Redundant Array of Independent Disks



(a) RAID 0 (non-redundant)

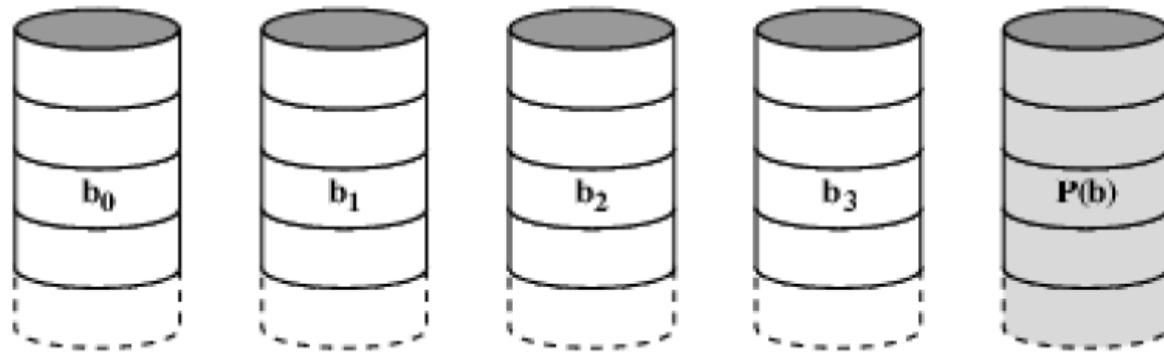


(b) RAID 1 (mirrored)

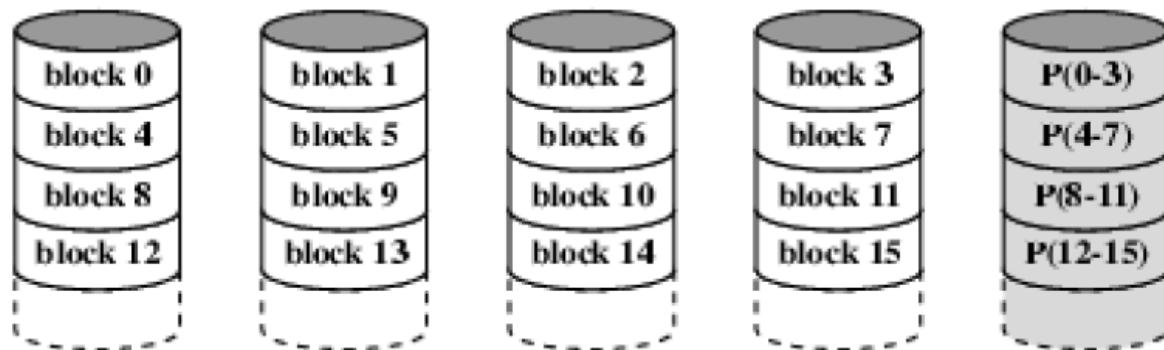


(c) RAID 2 (redundancy through Hamming code)

RAID 3 & 4

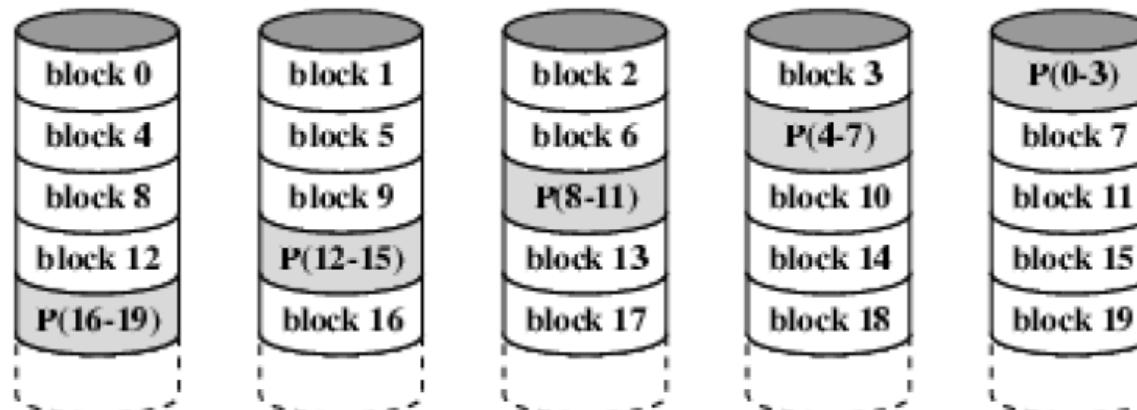


(d) RAID 3 (bit-interleaved parity)

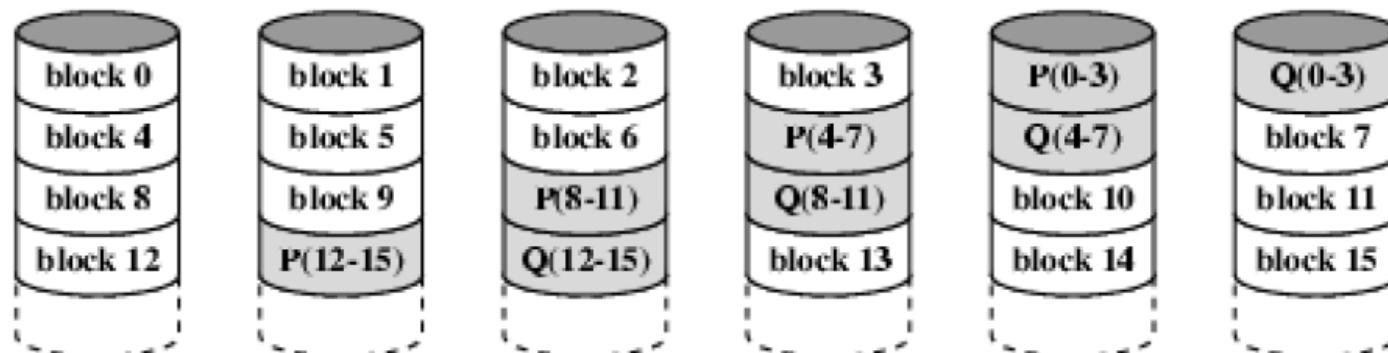


(e) RAID 4 (block-level parity)

RAID 5 & 6



(f) RAID 5 (block-level distributed parity)



(g) RAID 6 (dual redundancy)